


# ECE4514

## DIGITAL DESIGN 2

PERFORMANCE

CHAPTER 1





# GOING BEYOND “GETTING THE JOB DONE”

- High-Throughput Architectures
- Low-Latency Architectures
- Timing Optimizations
  - Adding Register Layers
  - Parallel Structures
  - Flatten Logic Structures
  - Register Balancing
  - Reorder Paths






# Throughput vs. Latency



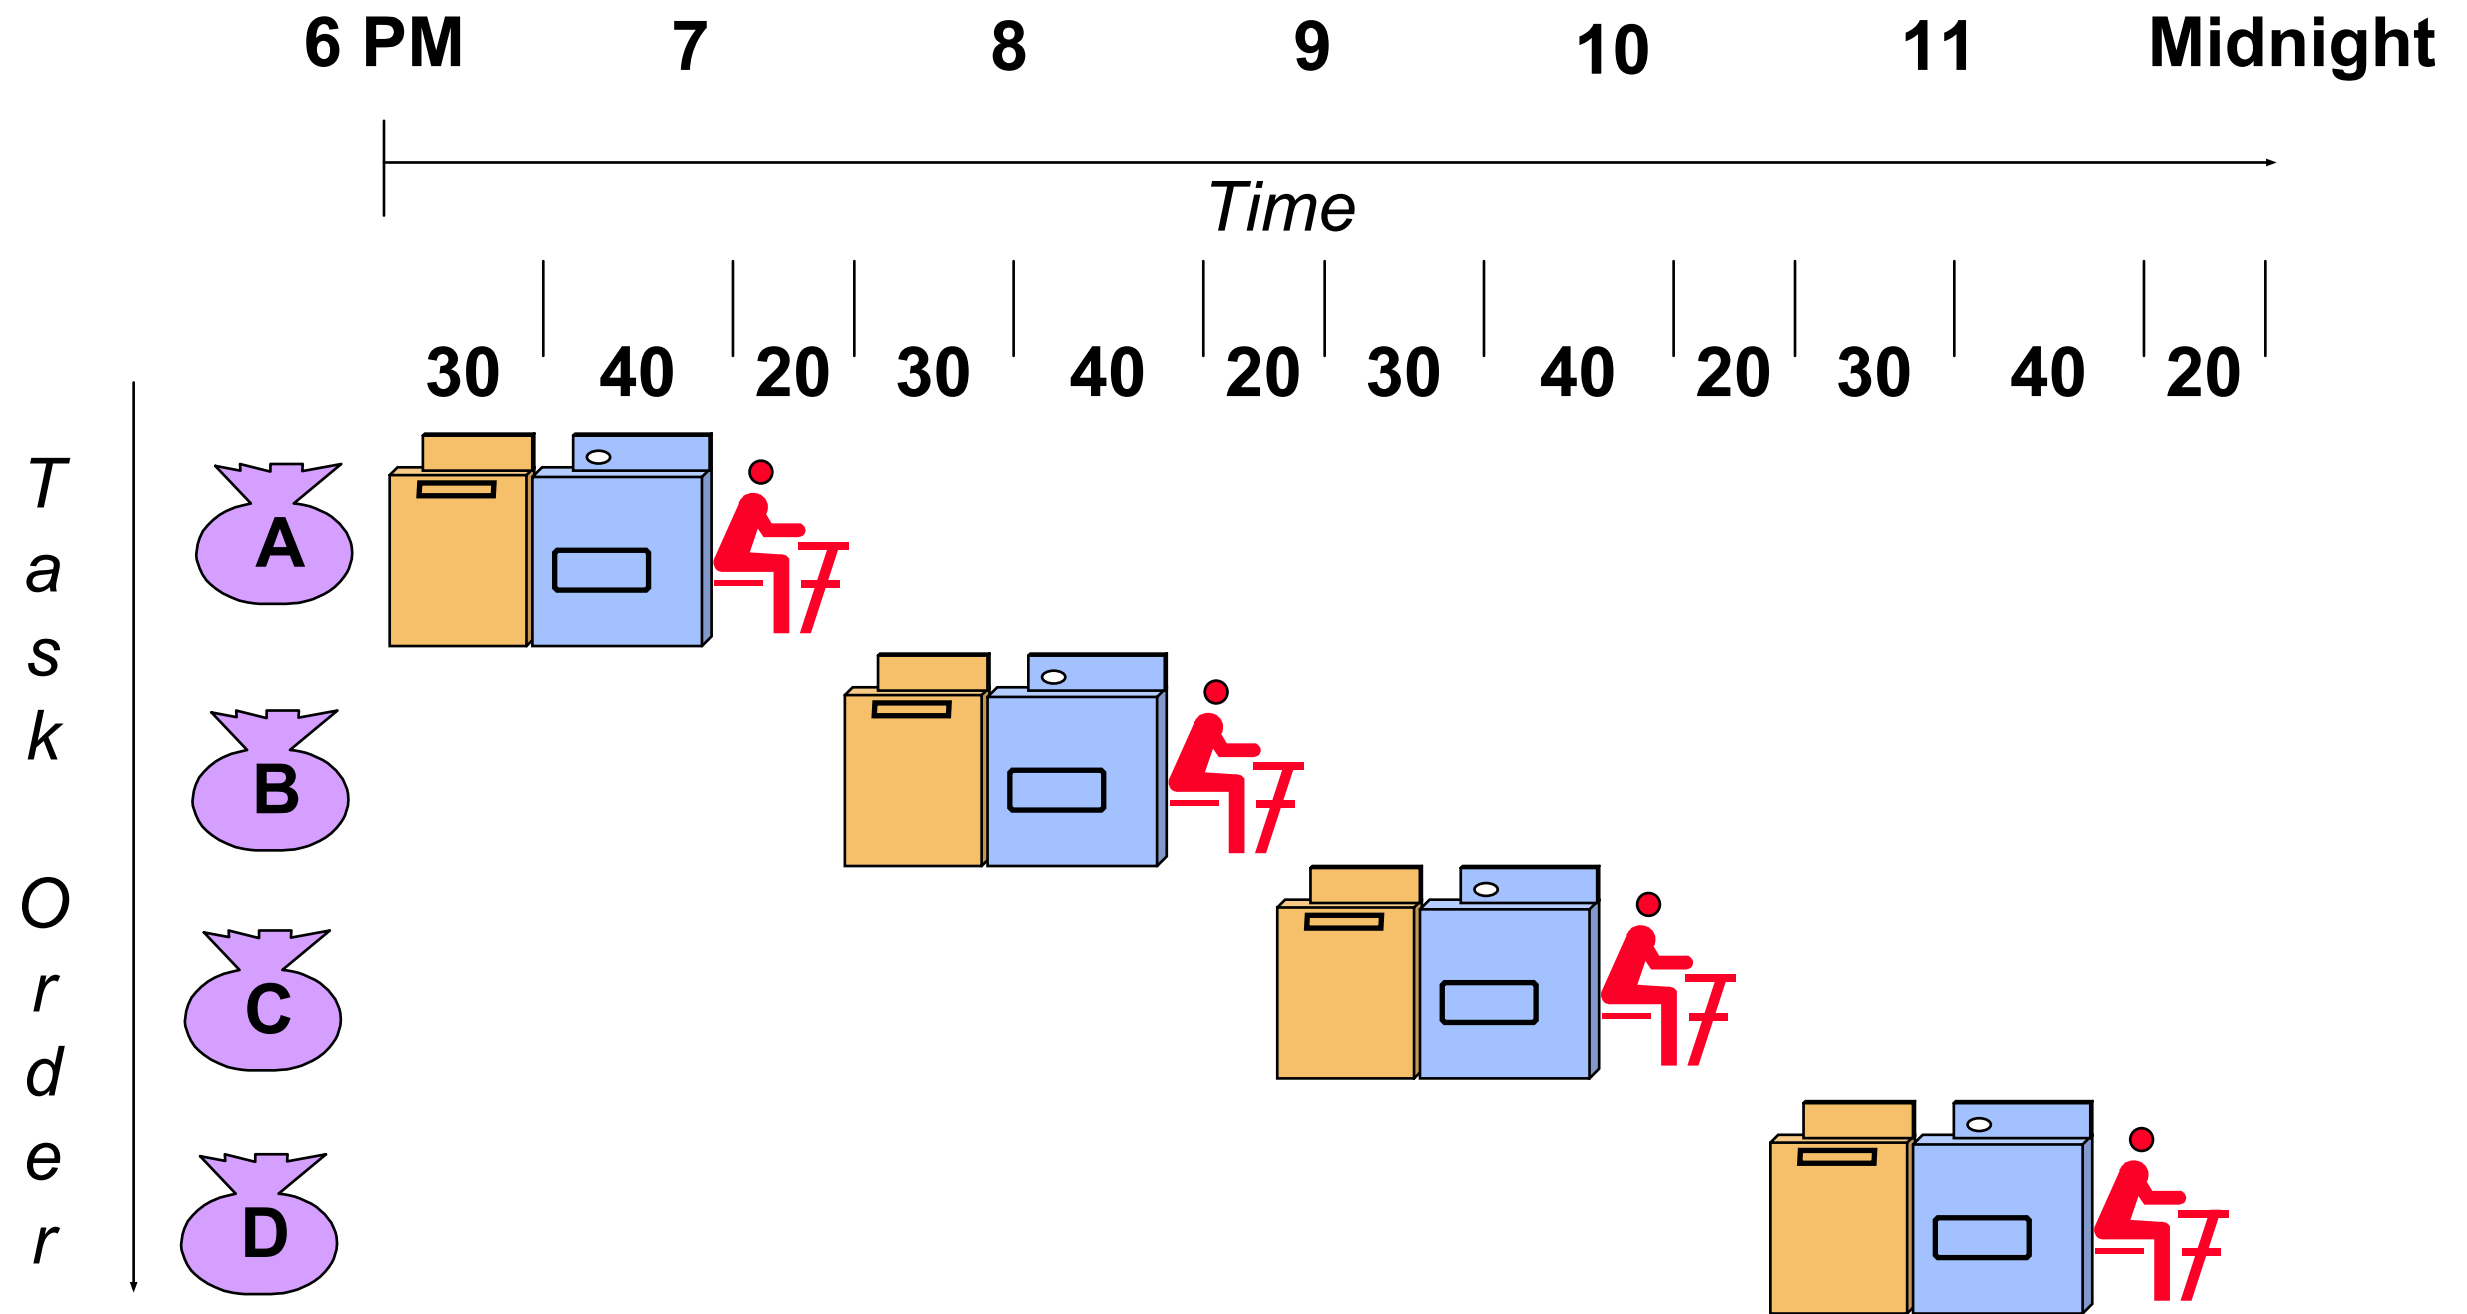
# Laundry Day

Ann, Brian, Cathy, Dave  
each have one load of clothes  
to wash, dry, and fold

- Washer takes 30 minutes
- Dryer takes 40 minutes
- Folding takes 20 minutes



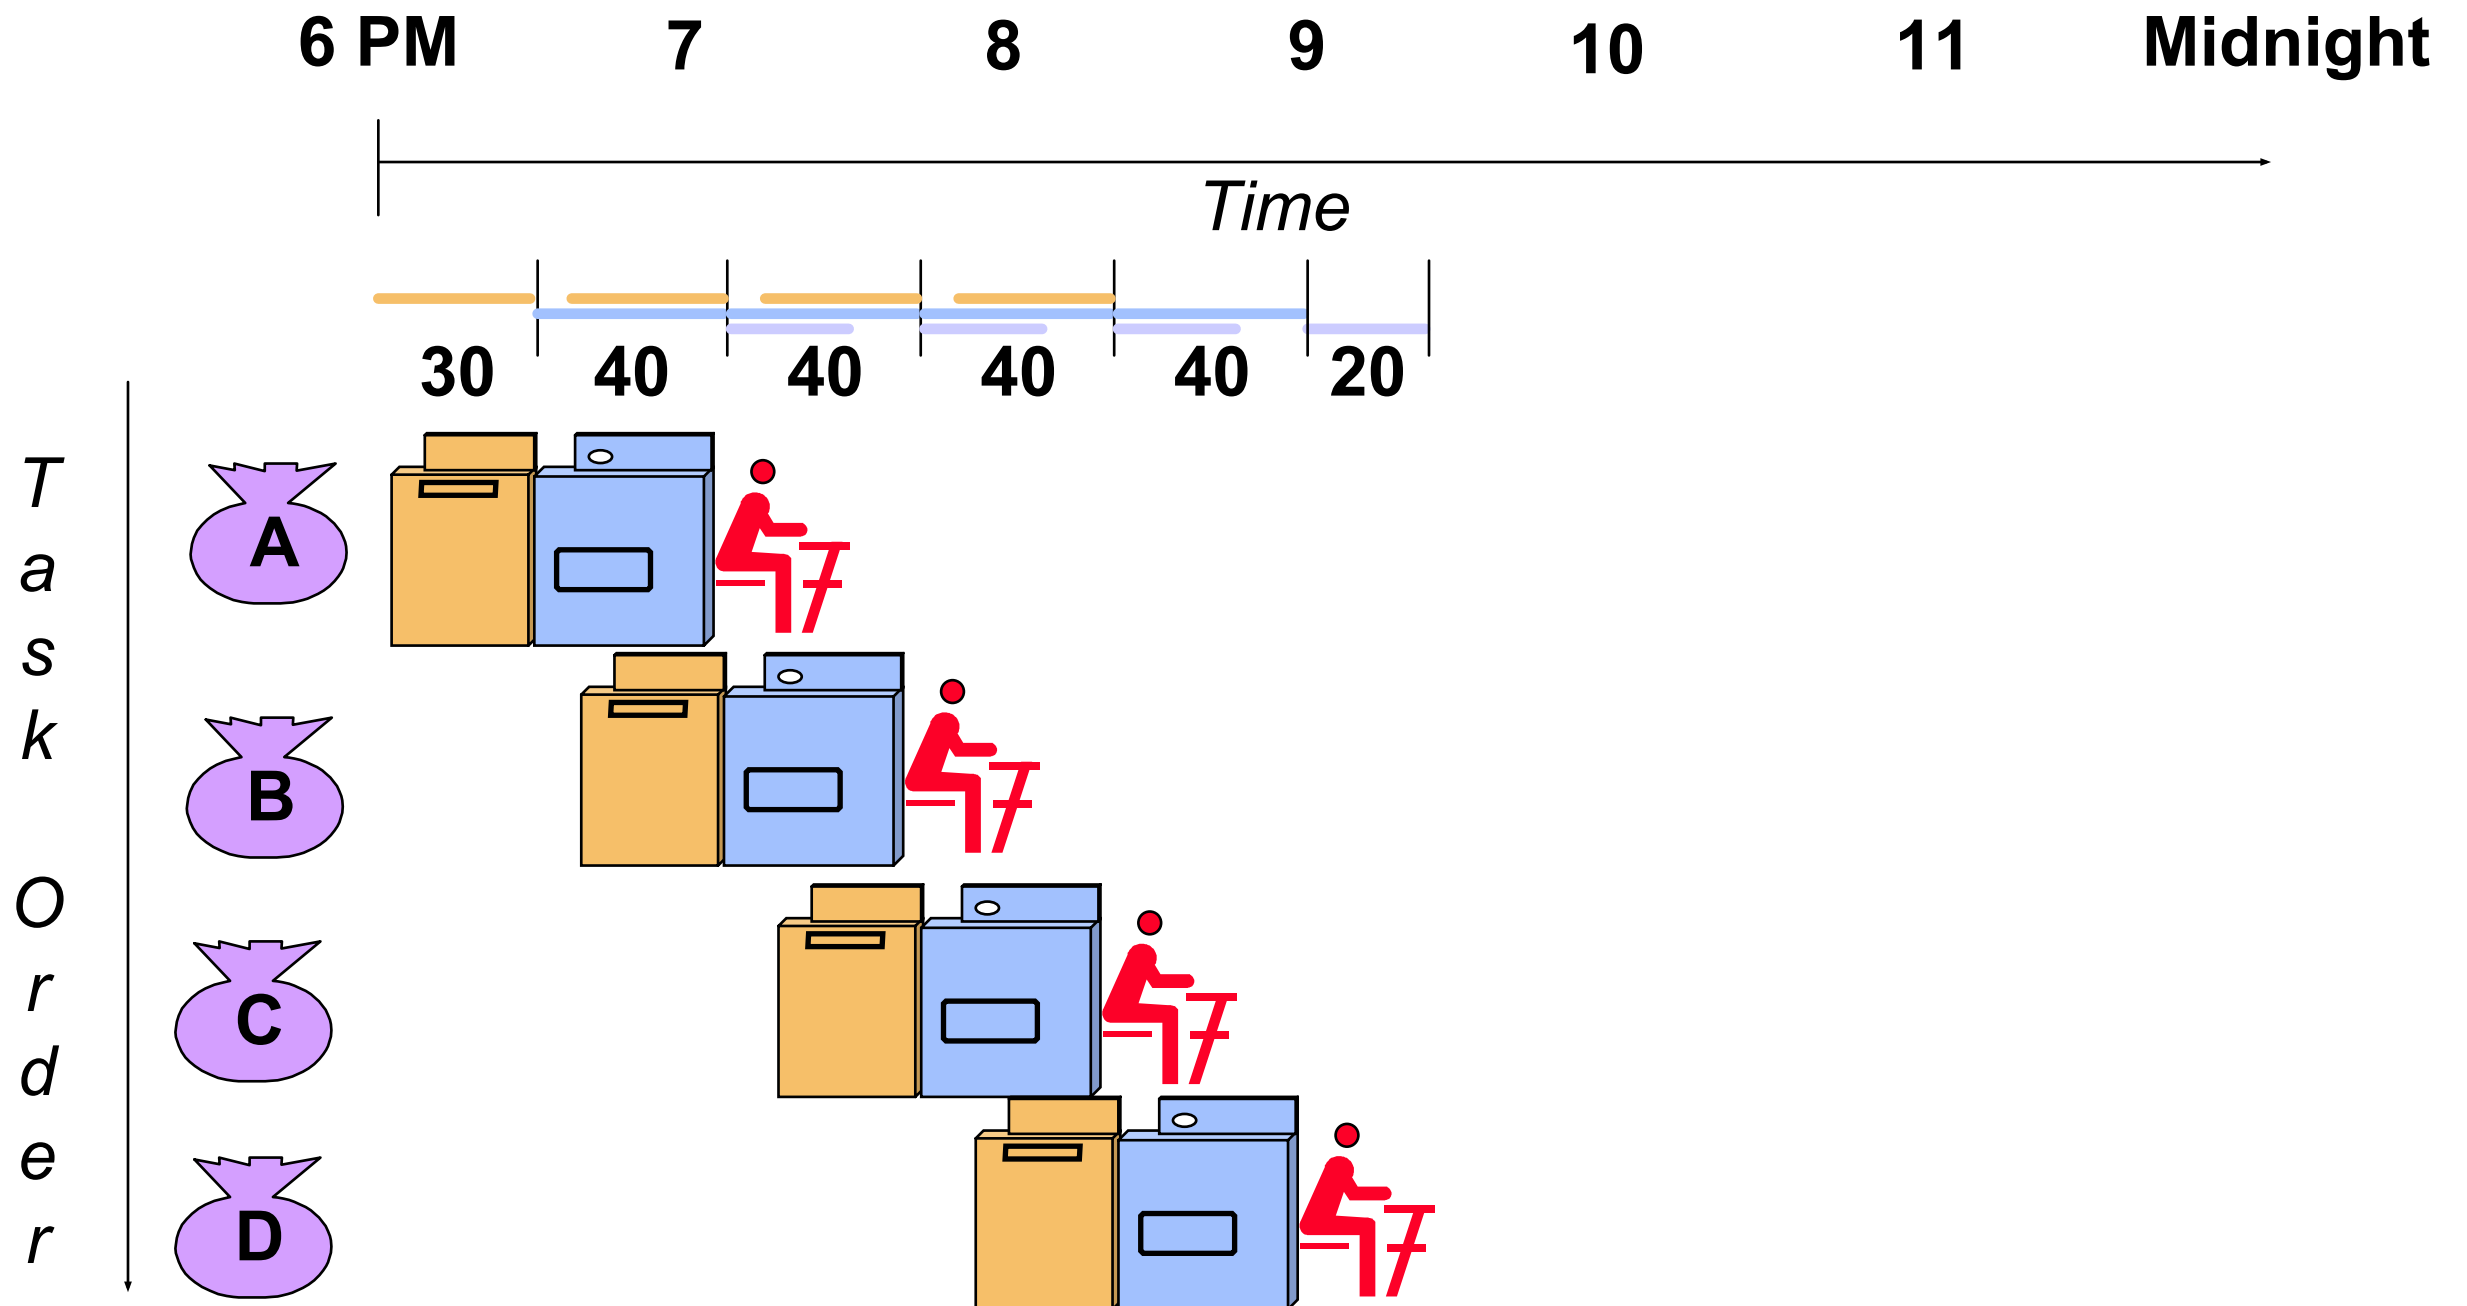
# Sequential Laundry



- Sequential laundry takes 6 hours for 4 loads
- With pipelining, how long would laundry take?

# Pipelined Laundry

## Start work ASAP



- Pipelined laundry takes 3.5 hours for 4 loads


# Key Definitions

*Pipelining* is a key implementation technique used to build fast processors. It allows the execution of multiple instructions to overlap in time.

A pipeline within a processor is similar to a car assembly line. Each assembly station is called a *pipe stage* or a *pipe segment*.

The **throughput** of an instruction pipeline is the measure of how often an instruction exits the pipeline.

# Pipeline Throughput and Latency



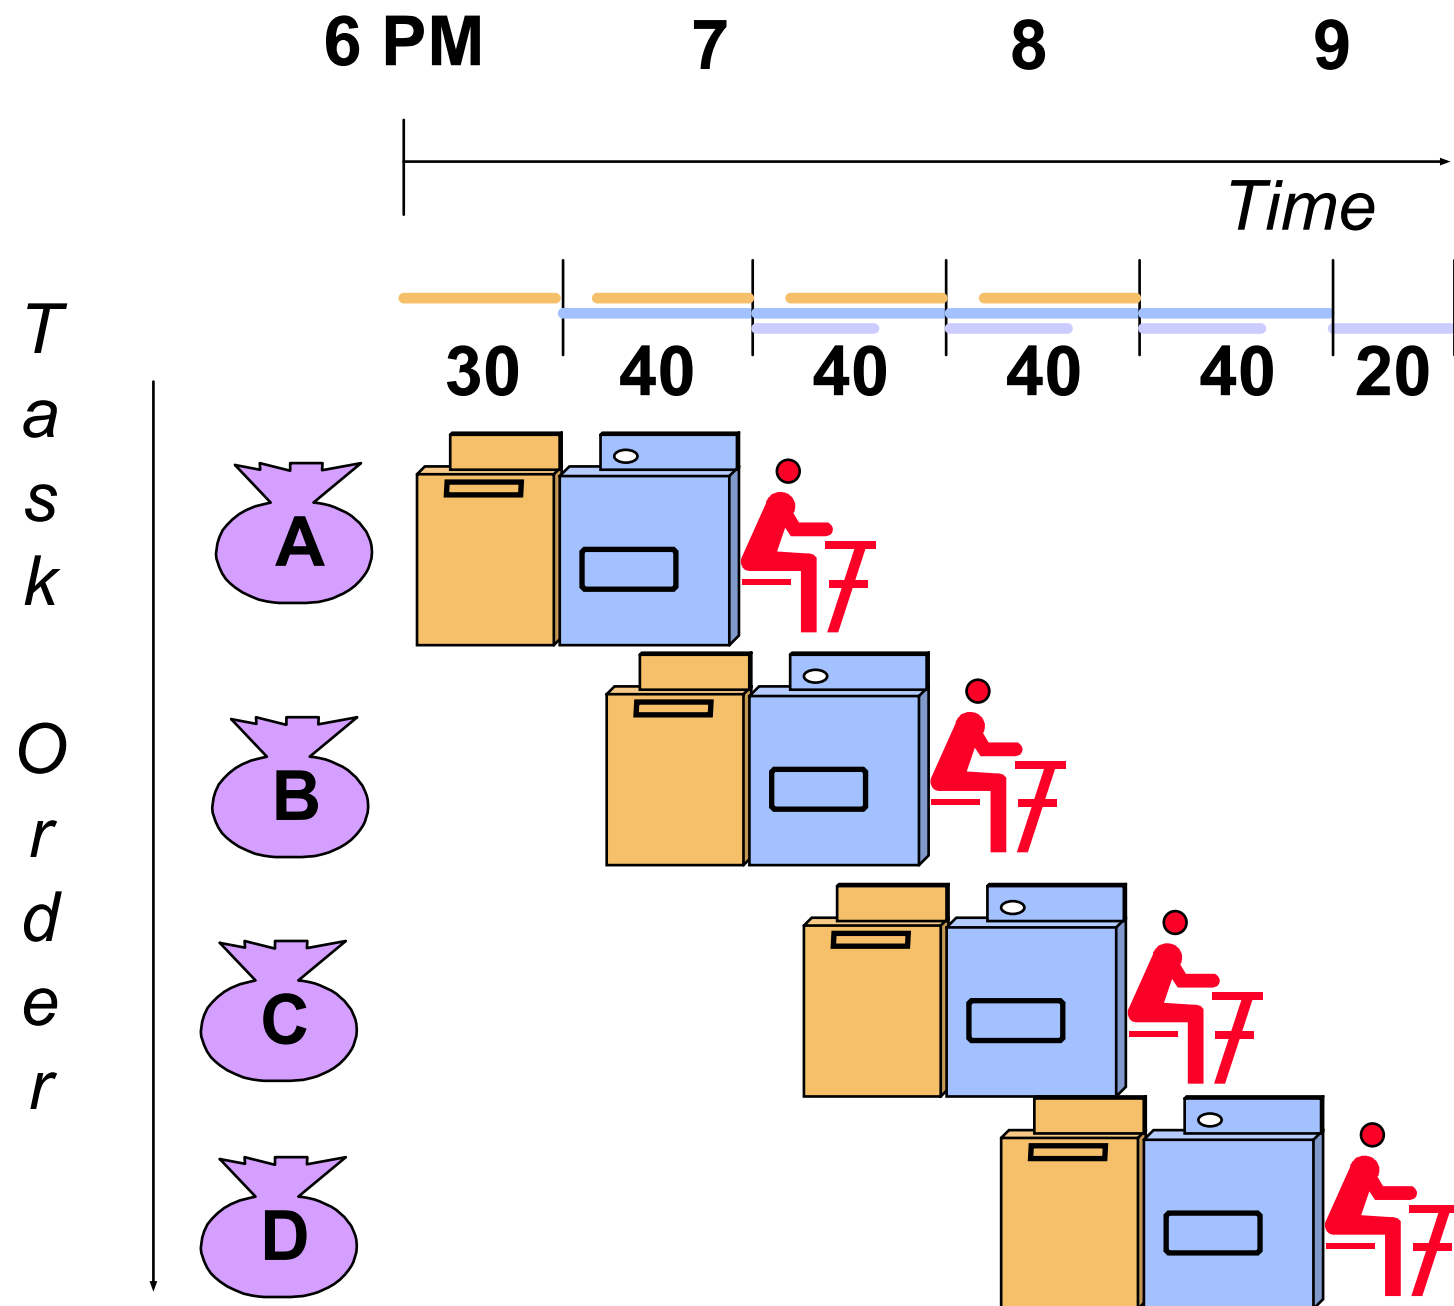
Consider the pipeline above with the indicated delays. We want to know what is the *pipeline throughput* and the *pipeline latency*.

Pipeline throughput: instructions completed per second.

Pipeline latency: how long does it take to execute a single instruction in the pipeline.



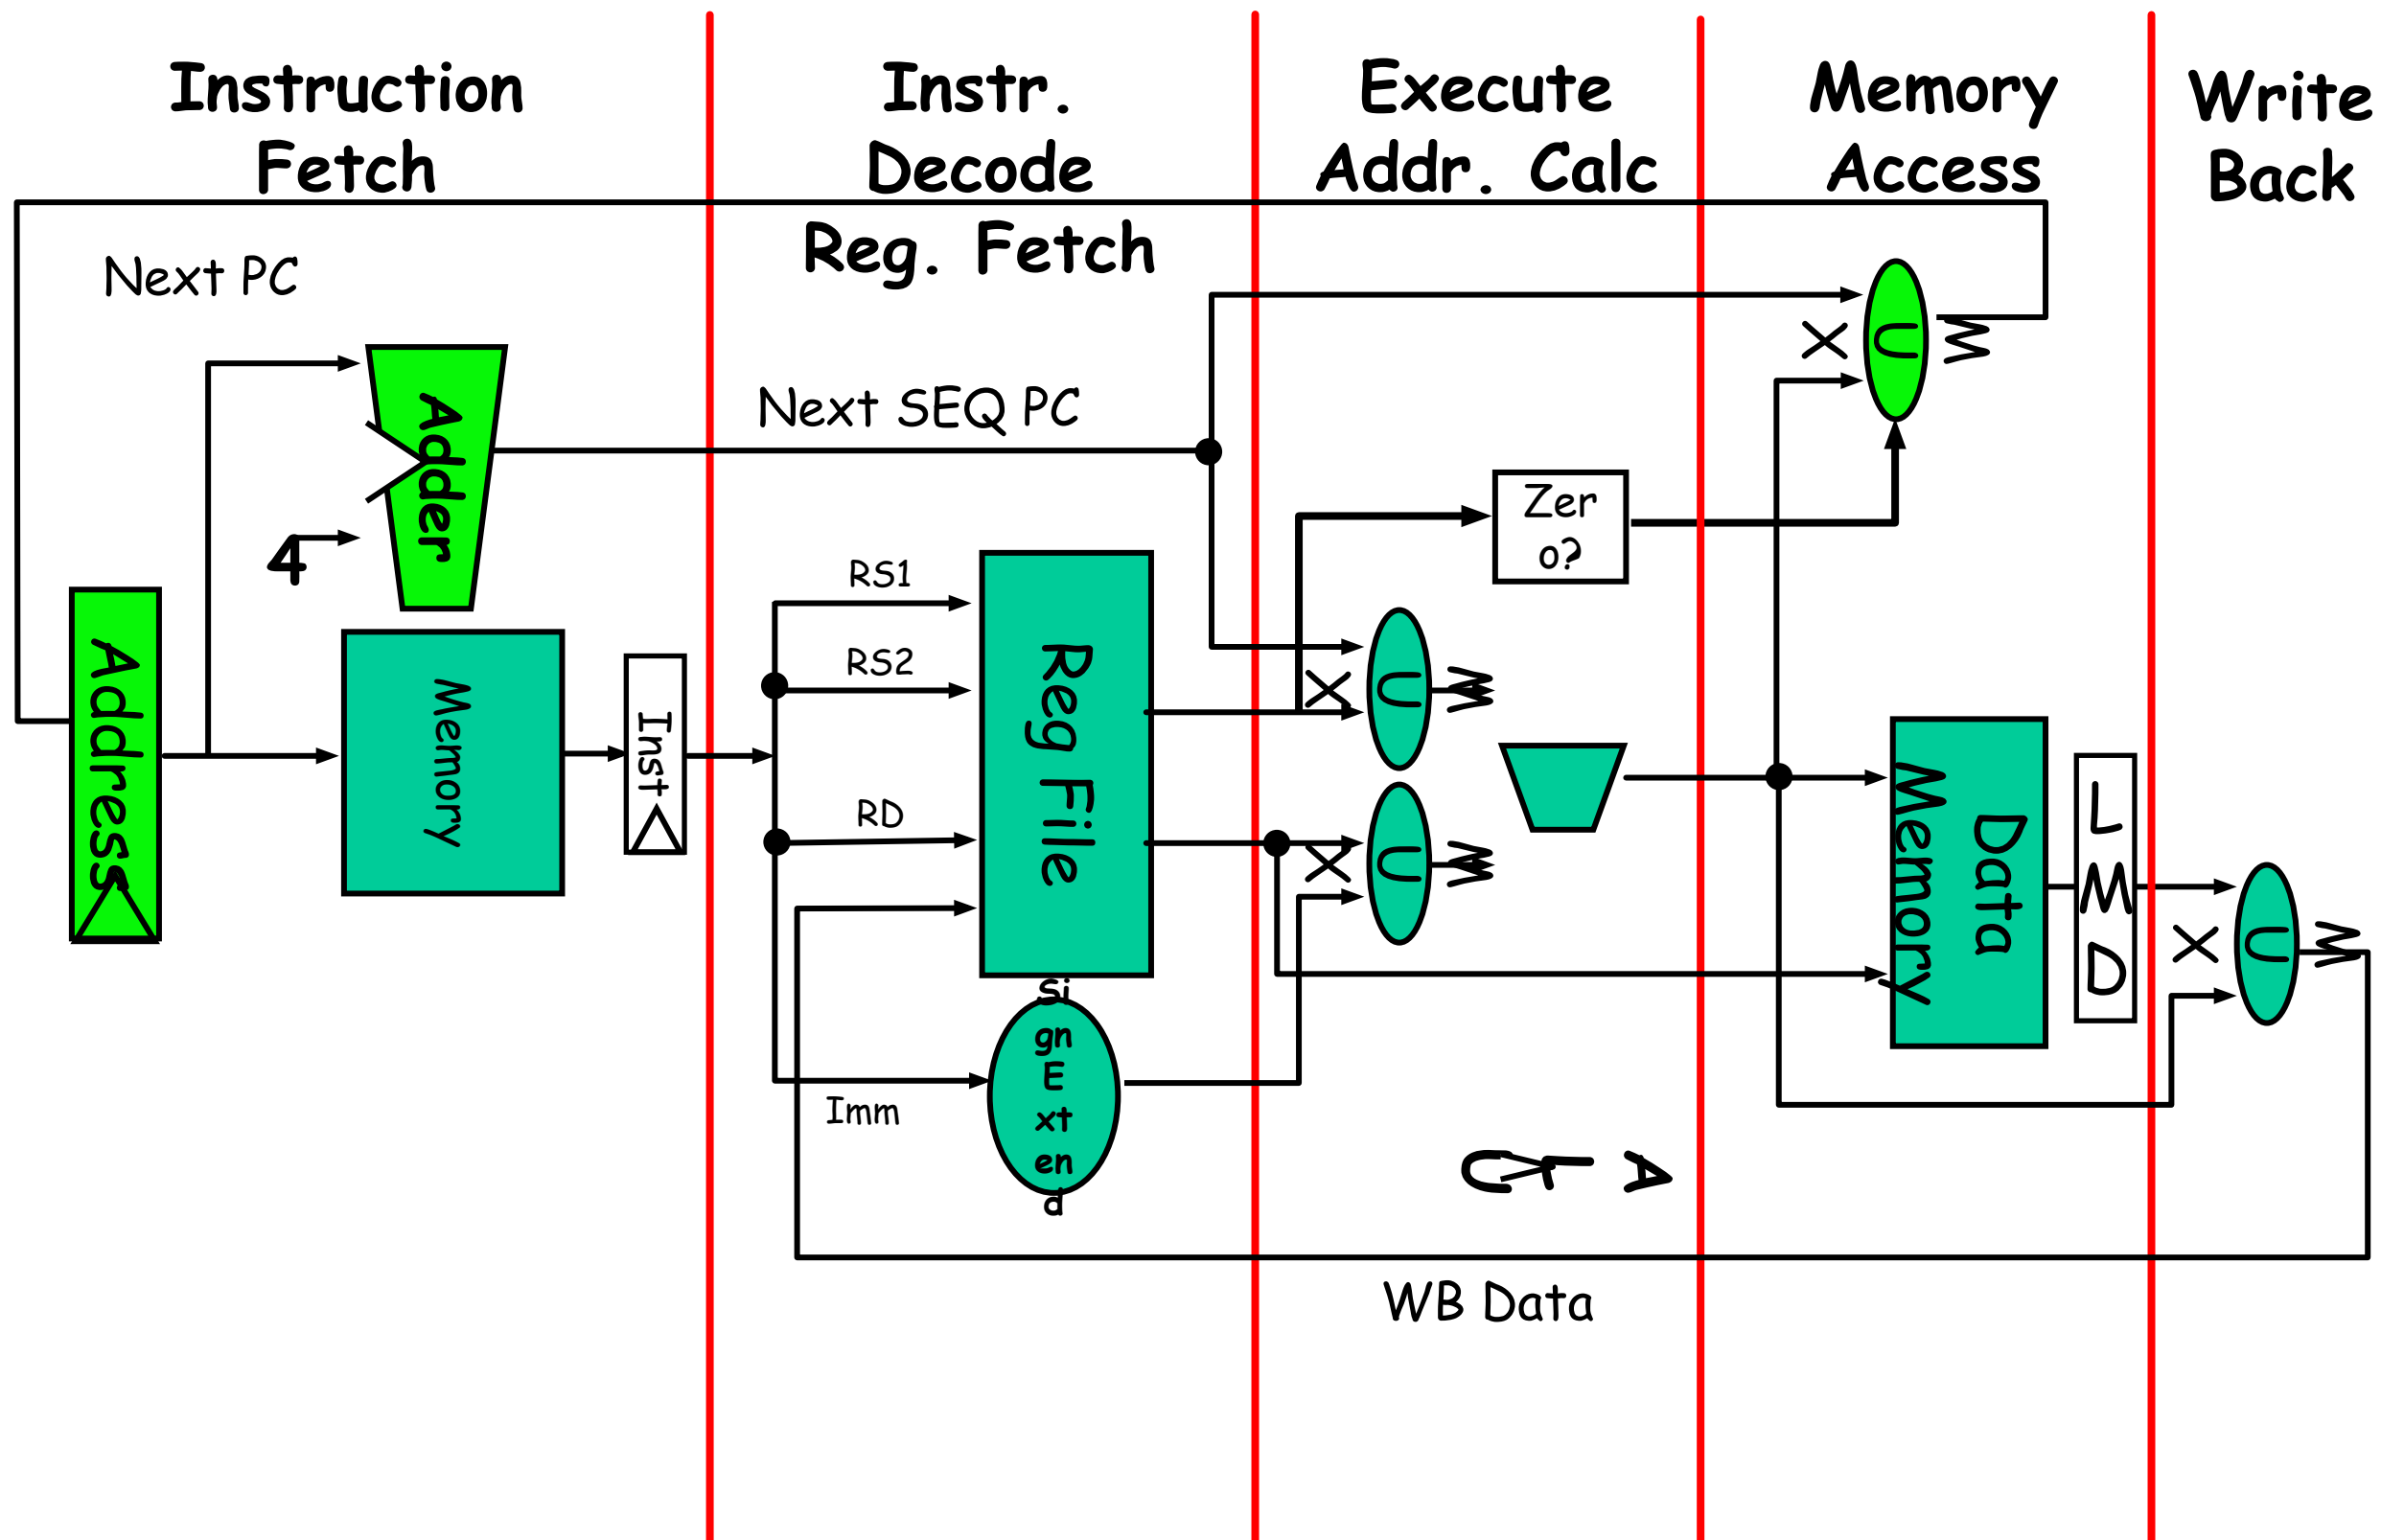
# Pipelining Lessons



- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Pipeline rate limited by **slowest** pipeline stage
- **Multiple** tasks operating simultaneously
- Potential speedup = **Number pipe stages**
- Unbalanced lengths of pipe stages reduces speedup
- Time to “**fill**” pipeline and time to “**drain**” it reduces speedup



# 5 Steps of MIPS Datapath





# Example 1: Design for low latency (parallelism)


$$\mathbf{X} = \mathbf{a} + \mathbf{b} + \mathbf{c} + \mathbf{d}$$

*All are 32-bit values*



# Example 1: Design for delay

- $X = a + b + c + d$



Delay = 1\*add + Reg

Latency = 2 cycles


Throughput =  $N$  bits/clock

# Example 2: Design for delay

```
x=0 ;
```

```
for (i=0; i<4; i++)
```

```
    x+= a[i]*b;
```



**Delay: 1\*Mul + 1 Add**

**Latency: 4 cycles**

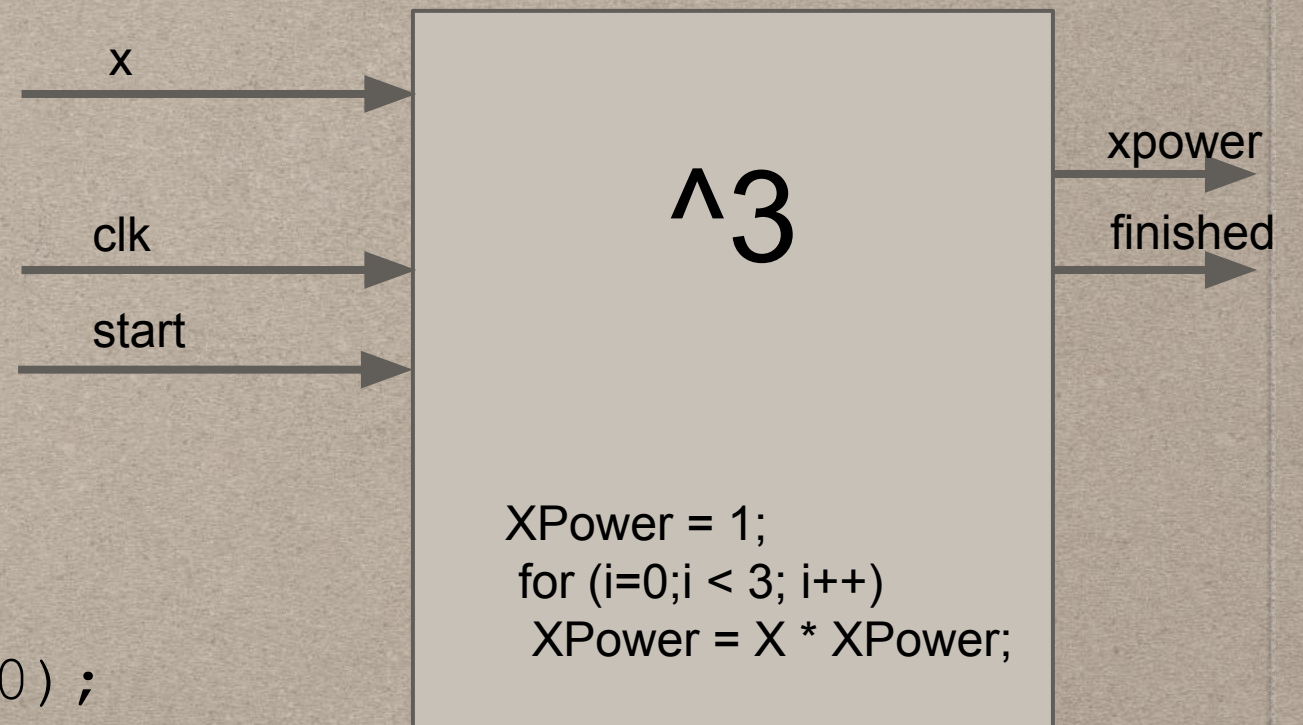
**Throughput:  $N$  bits/4 cycles**



# HIGH-THROUGHPUT ARCHITECTURE --

## EXAMPLE 1

```
module power3(  
    output reg [7:0] xpower,  
    output finished,  
    input [7:0] x,  
    input clk, start);  
    reg [7:0] ncount;  
  
    assign finished = (ncount == 0);  
    always@(posedge clk)  
        if(start) begin  
            xpower <= x;  
            ncount <= 2;  
        end  
        else if(!finished) begin  
            ncount <= ncount - 1;  
            xpower <= xpower * x;  
        end  
    end  
endmodule
```



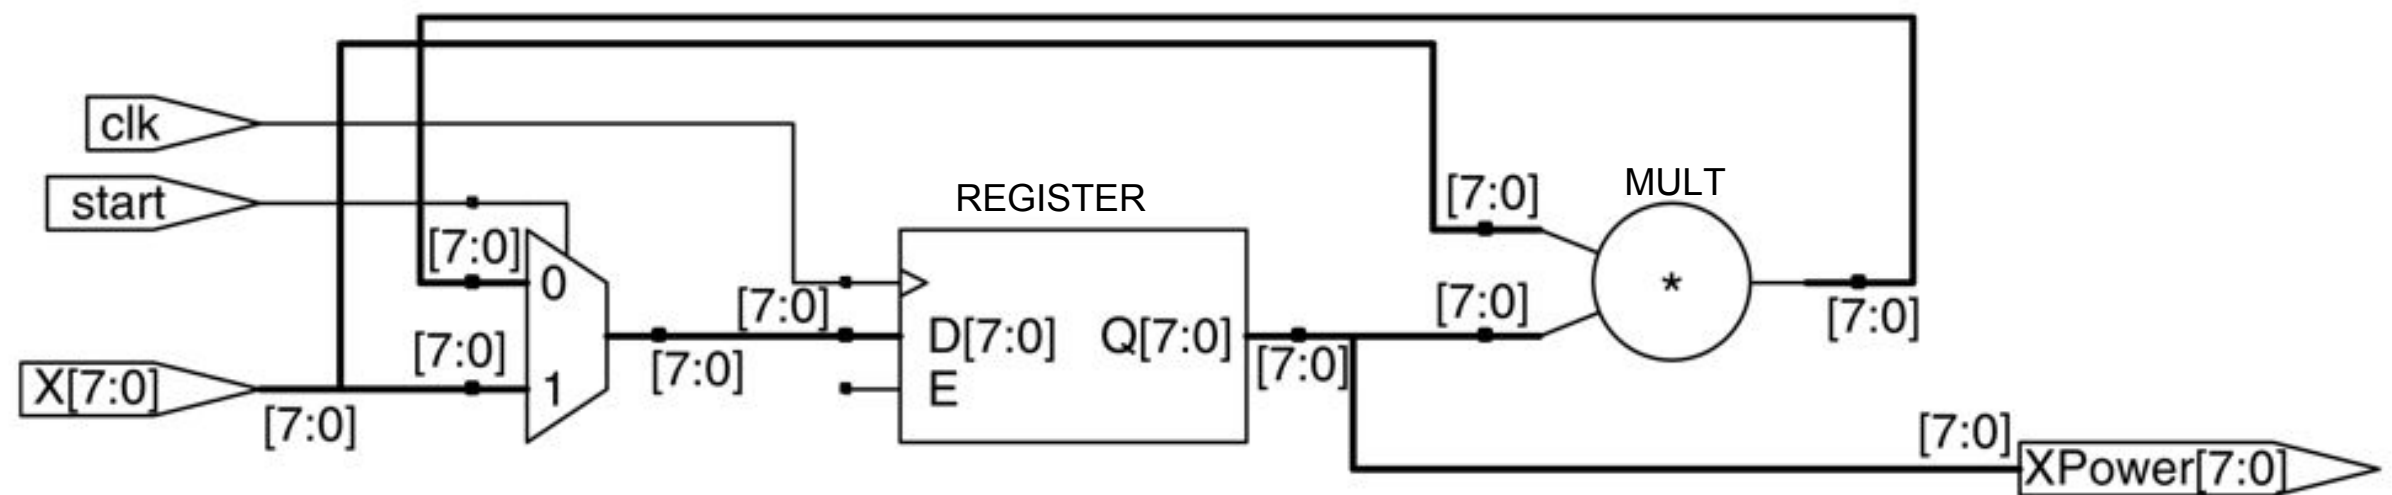
THROUGHPUT =

LATENCY =

CYCLE TIME =



# HIGH-THROUGHPUT ARCHITECTURE -- EXAMPLE 1



THROUGHPUT =

LATENCY =

CRITICAL PATH =

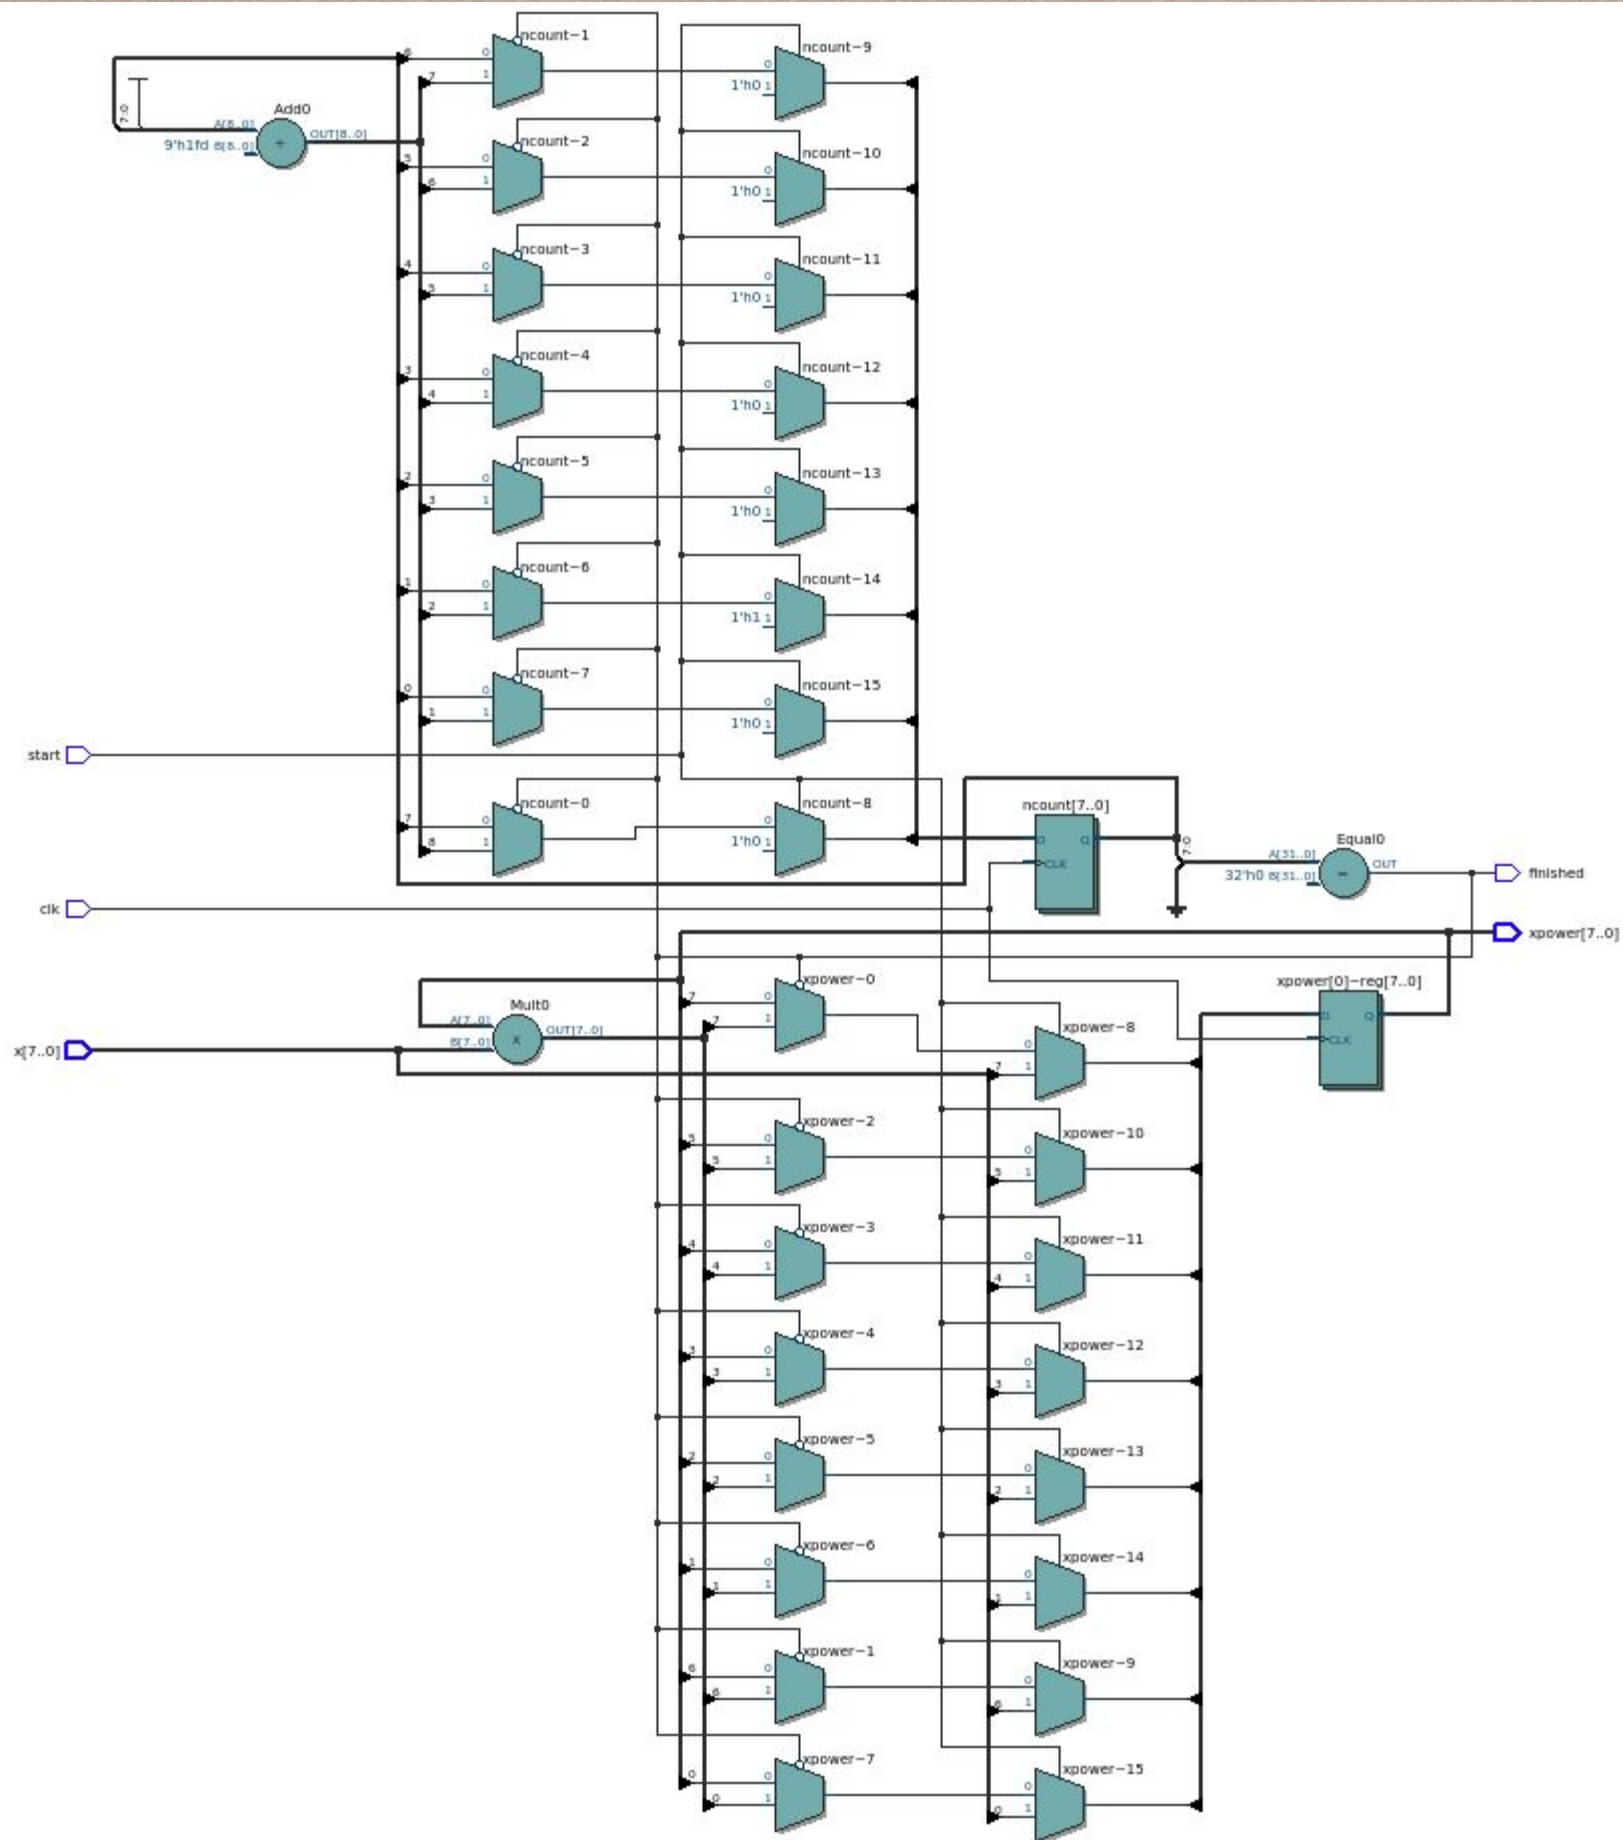


# ALTERA *FLOW* SUMMARY REPORT

Top-level Entity Name	power3
Family	Cyclone V
Device	5CSEMA5F31C6
Timing Models	Final
Logic utilization (in ALMs)	14 / 32,070 ( < 1 % )
Total registers	16
Total pins	19 / 457 ( 4 % )
Total virtual pins	0
Total block memory bits	0 / 4,065,280 ( 0 % )
Total DSP Blocks	1 / 87 ( 1 % )
Total HSSI RX PCSs	0
Total HSSI PMA RX Deserializers	0
Total HSSI TX PCSs	0
Total HSSI PMA TX Serializers	0
Total PLLs	0 / 6 ( 0 % )
Total DLLs	0 / 4 ( 0 % )

- [Adaptive Logic Modules \(ALMs\)](#)
- [DSP block 9-bit elements](#) -- digital signal processing block 9-bit element
- **HSSI RX PCSs** -- high speed serial interface physical coding sub-layer receiver channels
- **HSSI PMA RX Deserializers** -- high speed serial interface physical media attachment receiver channels
- **HSSI TX PCSs** -- high speed serial interface physical coding sub-layer transmitter channels
- **HSSI PMA TX Serializers** -- high speed serial interface physical media attachment transmitter channels.
- [PLLs](#)
- [DLLs](#)





# Altera RTL Mapping of Power3



# HIGH-THROUGHPUT ARCHITECTURE -- EXAMPLE 2

```
module power3b(  
    output reg [7:0] XPower,  
    input          clk,  
    input          [7:0] X  
);  
    reg          [7:0] XPower1, XPower2;  
    reg          [7:0] X1, X2;  
    always @(posedge clk) begin  
        // Pipeline stage 1  
        X1          <= X;  
        XPower1 <= X;  
        // Pipeline stage 2  
        X2          <= X1;  
        XPower2 <= XPower1 * X1;  
        // Pipeline stage 3  
        XPower <= XPower2 * X2;  
    end  
endmodule
```



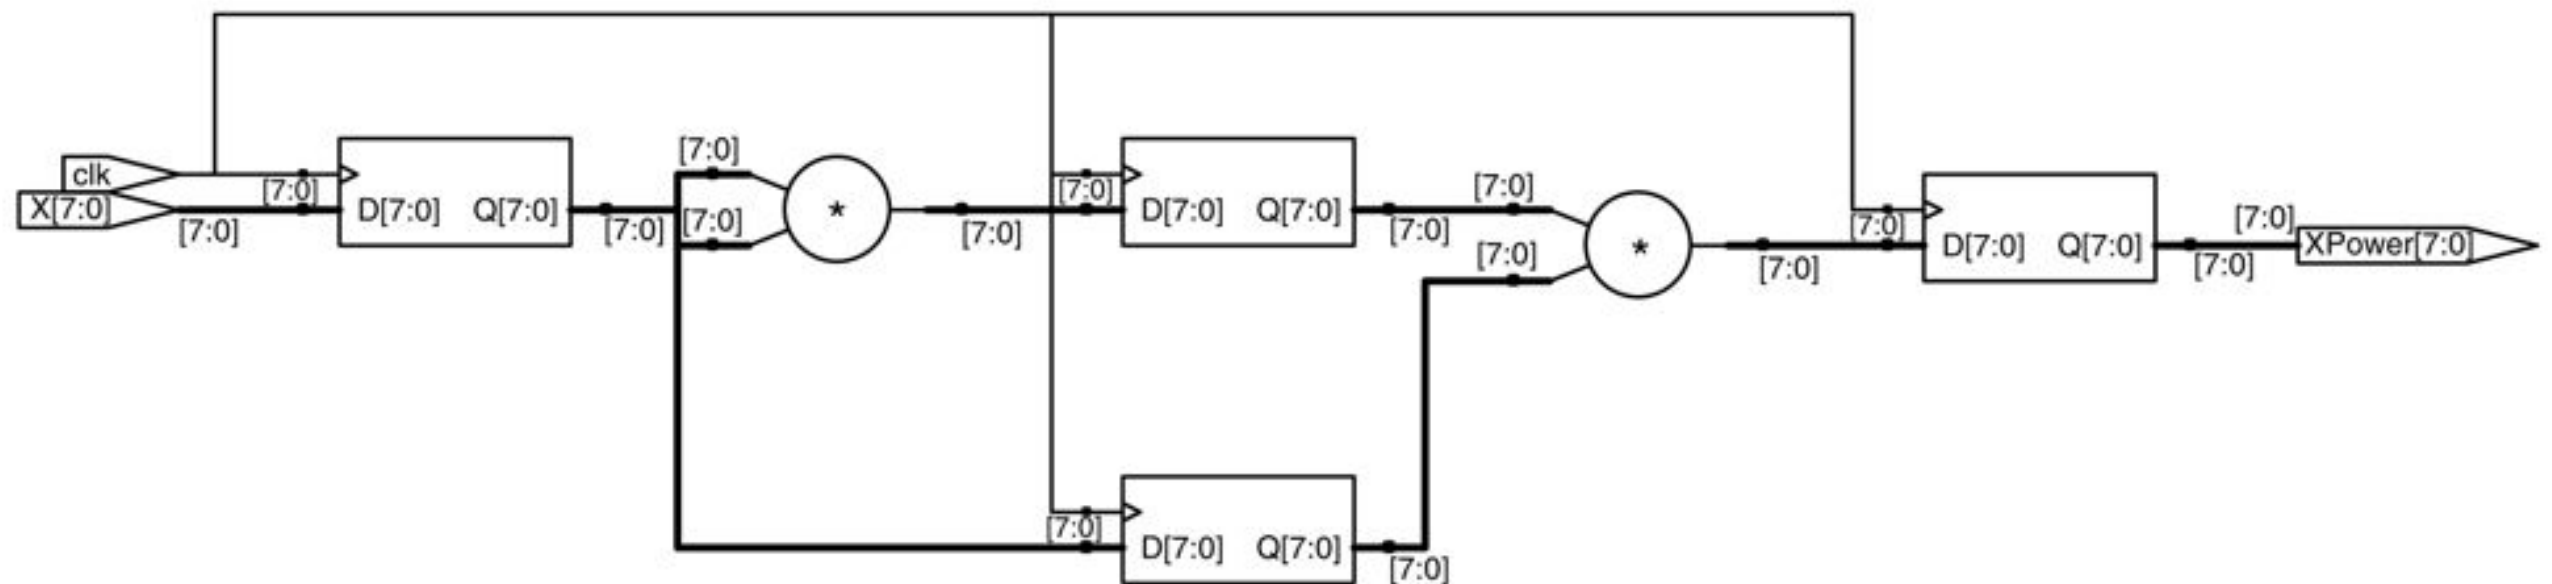
THROUGHPUT =

LATENCY =

CRITICAL PATH =



# HIGH-THROUGHPUT ARCHITECTURE -- EXAMPLE 2



THROUGHPUT =

LATENCY =

CYCLE TIME =



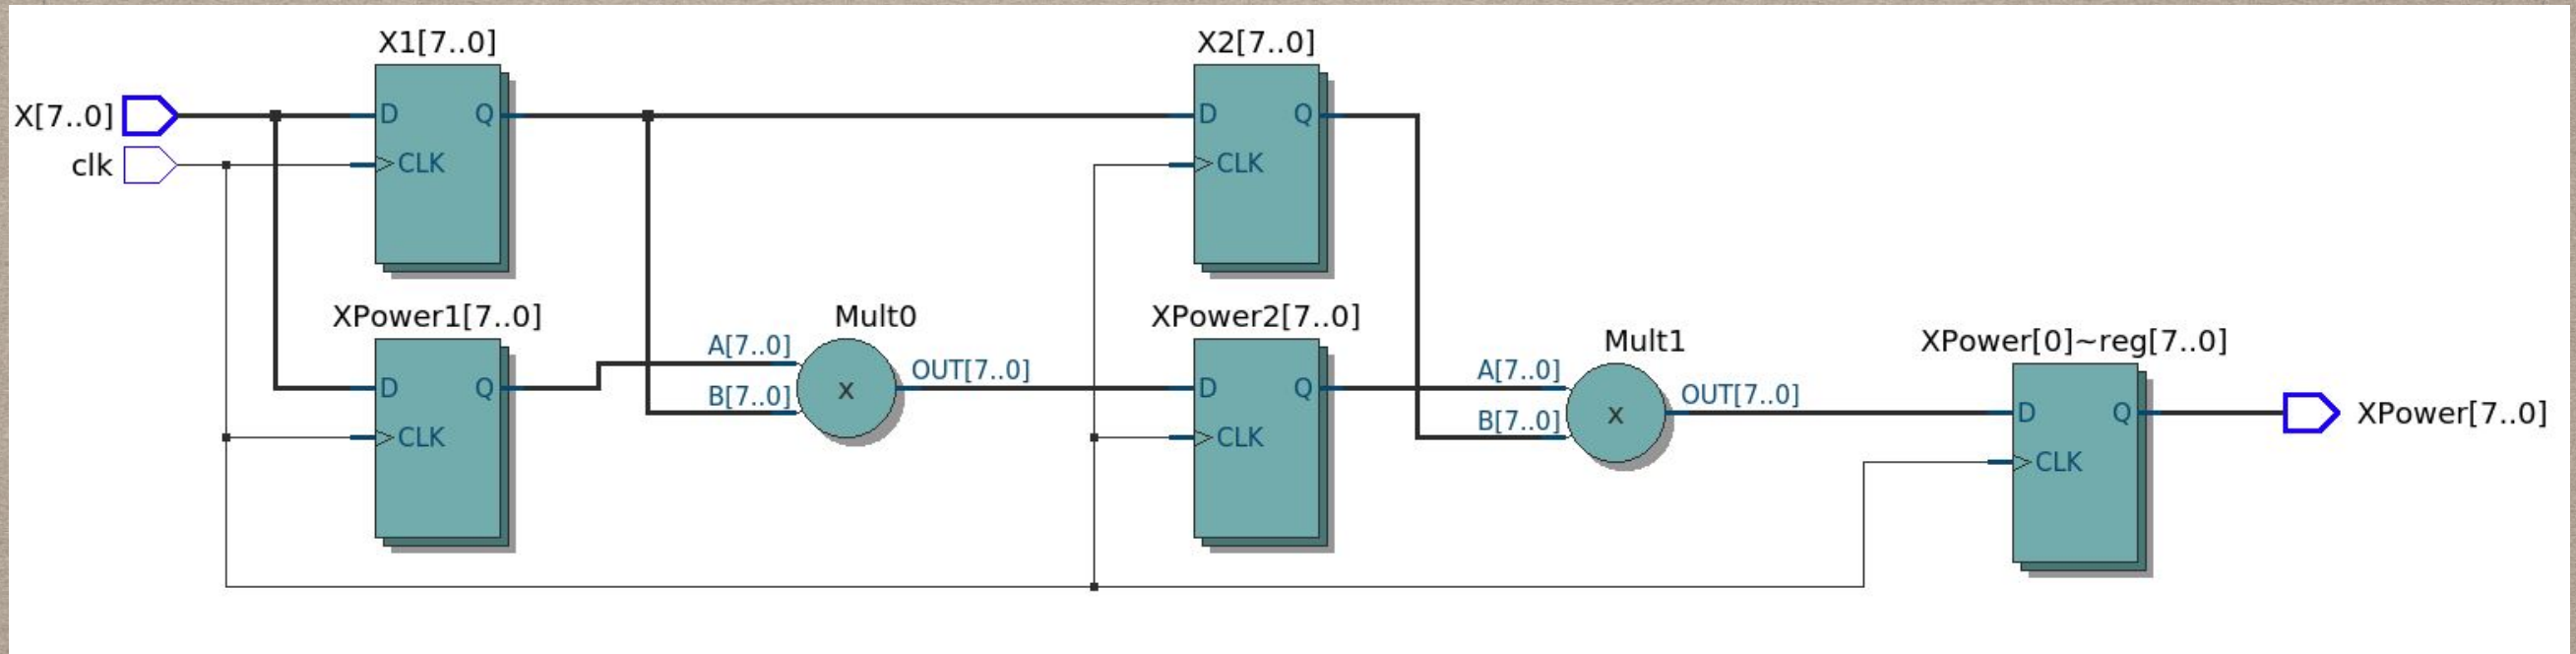
# ALTERA *FLOW SUMMARY* REPORT

Top-level Entity Name	power3b
Family	Cyclone V
Device	5CSEMA5F31C6
Timing Models	Final
Logic utilization (in ALMs)	3 / 32,070 ( < 1 % )
Total registers	8
Total pins	17 / 457 ( 4 % )
Total virtual pins	0
Total block memory bits	0 / 4,065,280 ( 0 % )
Total DSP Blocks	2 / 87 ( 2 % )
Total HSSI RX PCSs	0
Total HSSI PMA RX Deserializers	0
Total HSSI TX PCSs	0
Total HSSI PMA TX Serializers	0
Total PLLs	0 / 6 ( 0 % )
Total DLLs	0 / 4 ( 0 % )





# ALTERA MAPPING OF POWER3B





# LOW LATENCY ARCHITECTURE -- EXAMPLE 3

```
module power3(  
    output [7:0] XPower,  
    input  [7:0] X  
);  
    reg    [7:0] XPower1, XPower2;  
    reg    [7:0] X1, X2;  
    assign XPower = XPower2 * X2;  
    always @*  
    begin  
        X1 =X; XPower1 = X;  
    end  
    always @* begin  
        X2      = X1;  
        XPower2 = XPower1*X1;  
    end  
endmodule
```




THROUGHPUT =

LATENCY =

CRITICAL PATH =



# HIGH-THROUGHPUT ARCHITECTURE -- EXAMPLE 3



THROUGHPUT =


LATENCY =

CYCLE TIME =



# ALTERA *FLOW* SUMMARY REPORT

Top-level Entity Name	power3c
Family	Cyclone V
Device	5CSEMA5F31C6
Timing Models	Final
Logic utilization (in ALMs)	1 / 32,070 ( < 1 % )
Total registers	0
Total pins	16 / 457 ( 4 % )
Total virtual pins	0
Total block memory bits	0 / 4,065,280 ( 0 % )
Total DSP Blocks	2 / 87 ( 2 % )
Total HSSI RX PCSs	0
Total HSSI PMA RX Deserializers	0
Total HSSI TX PCSs	0
Total HSSI PMA TX Serializers	0
Total PLLs	0 / 6 ( 0 % )
Total DLLs	0 / 4 ( 0 % )





# ALTERA MAPPING OF POWER3C






# RESULTS

DESIGN	FMAX	ALM	DSP	REGS	THRU-PUT
power3a (1st example)	217 MHz	14	1	16	578 Mbps
power3b (2nd example)	257 MHz	3	2	8	2056 Mbps
power3c (no clock)	---	1	2	0	----
power3d (3rd example)	178 MHz	1	2	0	1424 Mbps
power3e (2nd redone with piped mults)	469 MHz	5	2	16	3752 Mbps



# TIMING


$$F_{\max} = \frac{1}{T_{\text{clk-q}} + T_{\text{logic}} + T_{\text{routing}} + T_{\text{setup}} - T_{\text{skew}}}$$





# TIMING -- EXAMPLE: FINITE IMPULSE FILTER (FIR)

```
module fir(  
    output [7:0] Y,  
    input  [7:0] A, B, C, X,  
    input          clk,  
  
    input          validsample);  
    reg [7:0] X1, X2, Y;  
    always @(posedge clk)  
        if(validsample) begin  
            X1 <= X;  
            X2 <= X1;  
            Y <= A* X+B* X1+C* X2;  
        end  
endmodule
```



THROUGHPUT =

LATENCY =

CRITICAL PATH =



# TIMING -- ADD REGISTER LAYERS

```

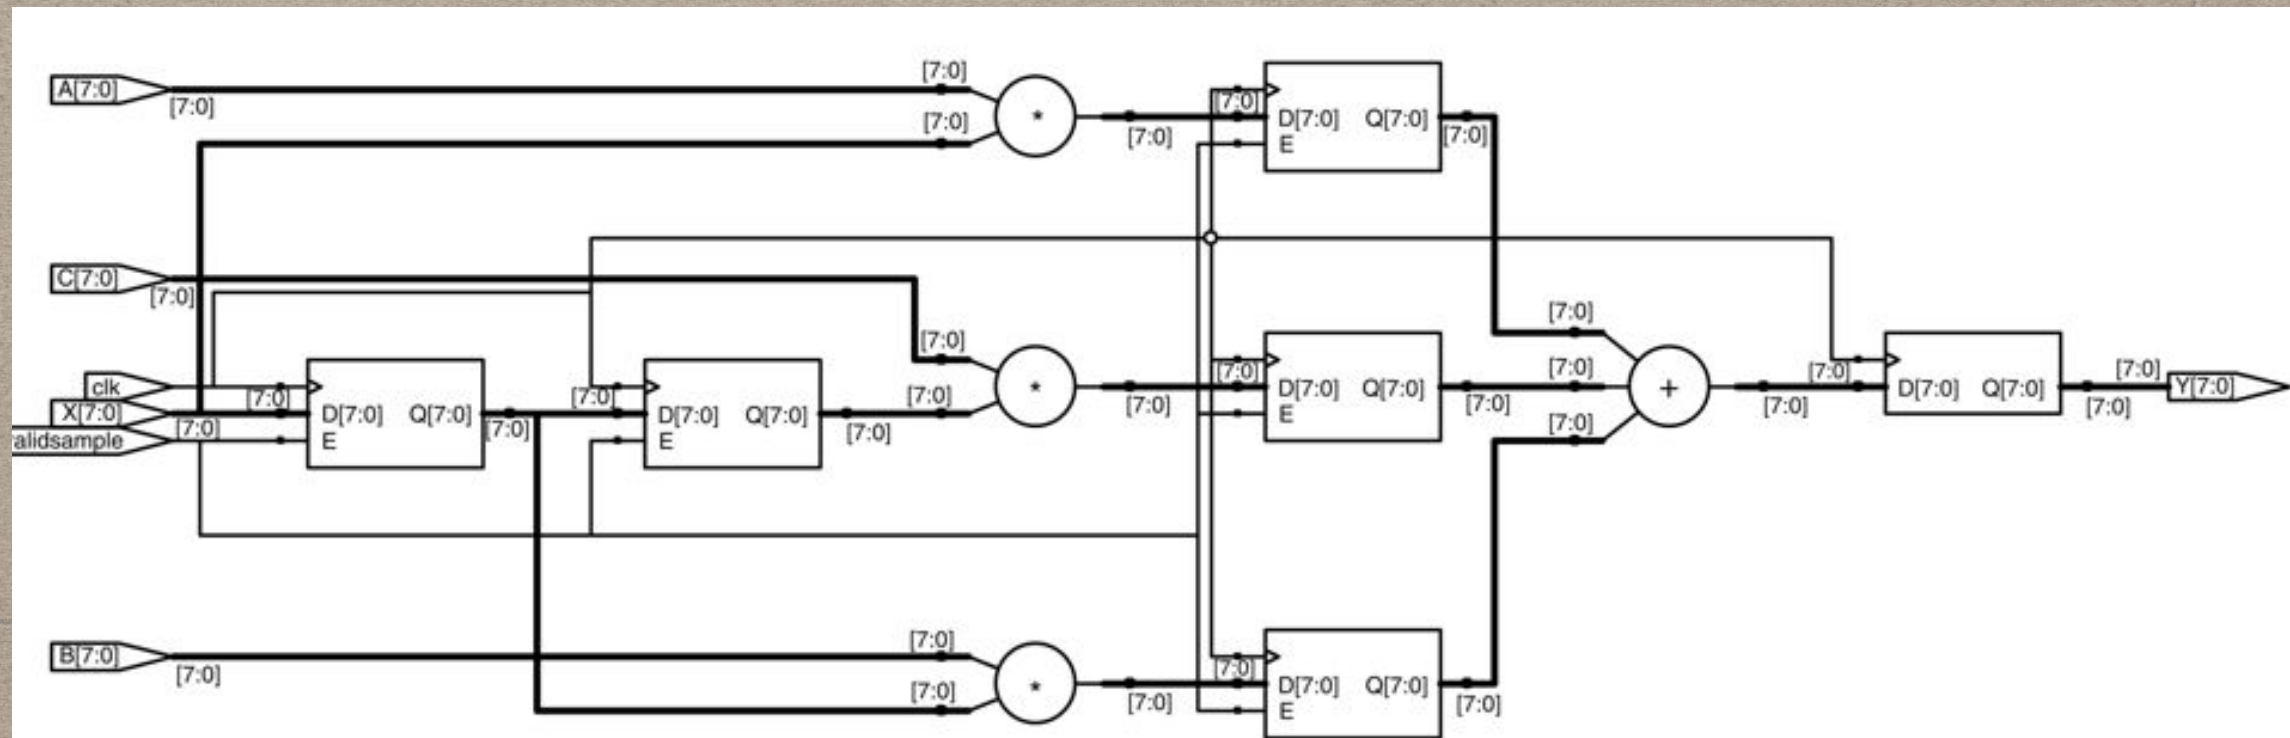
module fir( ----- );
    reg      [7:0] X1, X2, Y;
    reg      [7:0] prod1, prod2, prod3;
    always @ (posedge clk) begin
        if(validsample) begin
            X1 <= X;
            X2 <= X1;
            prod1 <= A * X;
            prod2 <= B * X1;
            prod3 <= C * X2;
        end
        Y <= prod1 + prod2 + prod3;
    end
endmodule

```

THROUGHPUT =

LATENCY =

CRITICAL PATH =





# TIMING -- ADD MORE REGISTER LAYERS

```
module fir( ---- );
    reg      [7:0] X1, X2, Y;
    reg      [7:0] prod1, prod2, prod3, add1;
    always @ (posedge clk) begin
        if(validsample) begin
            X1 <= X;
            X2 <= X1;
            prod1 <= A * X;
            prod2 <= B * X1;
            prod3 <= C * X2;
            add1 <= prod1 + prod2;
        end
        Y <= add1 + prod3;
    end
endmodule
```

THROUGHPUT =

LATENCY =

CRITICAL PATH =



# TIMING -- PARALLEL STRUCTURES

```
module power3(  
    output [7:0] XPower,  
    input  [7:0] X,  
    input          clk);  
reg    [7:0] XPower1;  
// partial product registers  
reg    [3:0] XPower2_ppAA, XPower2_ppAB,  
        XPower2_ppBB;  
reg    [3:0] XPower3_ppAA, XPower3_ppAB,  
        XPower3_ppBB;  
reg    [7:0] X1, X2;  
wire   [7:0] XPower2;  
// nibbles for partial products (A is MS  
// nibble, B is LS nibble)  
wire   [3:0] XPower1_A = XPower1[7:4];  
wire   [3:0] XPower1_B = XPower1[3:0];  
wire   [3:0] X1_A      = X1[7:4];  
wire   [3:0] X1_B      = X1[3:0];  
wire   [3:0] XPower2_A = XPower2[7:4];  
wire   [3:0] XPower2_B = XPower2[3:0];  
wire   [3:0] X2_A      = X2[7:4];  
wire   [3:0] X2_B      = X2[3:0];
```

```
// assemble partial products  
assign XPower = (XPower2_ppAA << 8)+  
    (2*XPower2_ppAB << 4)+ XPower2_ppBB;  
assign XPower2 = (XPower3_ppAA << 8)+  
    (2*XPower3_ppAB << 4)+ XPower3_ppBB;  
always @(posedge clk) begin  
    // Pipeline stage 1  
    X1          <= X;  
    XPower1     <= X;  
    // Pipeline stage 2  
    X2          <= X1;  
    // create partial products  
    XPower2_ppAA <= XPower1_A * X1_A;  
    XPower2_ppAB <= XPower1_A * X1_B;  
    XPower2_ppBB <= XPower1_B * X1_B;  
    // Pipeline stage 3  
    // create partial products  
    XPower3_ppAA <= XPower2_A * X2_A;  
    XPower3_ppAB <= XPower2_A * X2_B;  
    XPower3_ppBB <= XPower2_B * X2_B;  
end  
endmodule
```

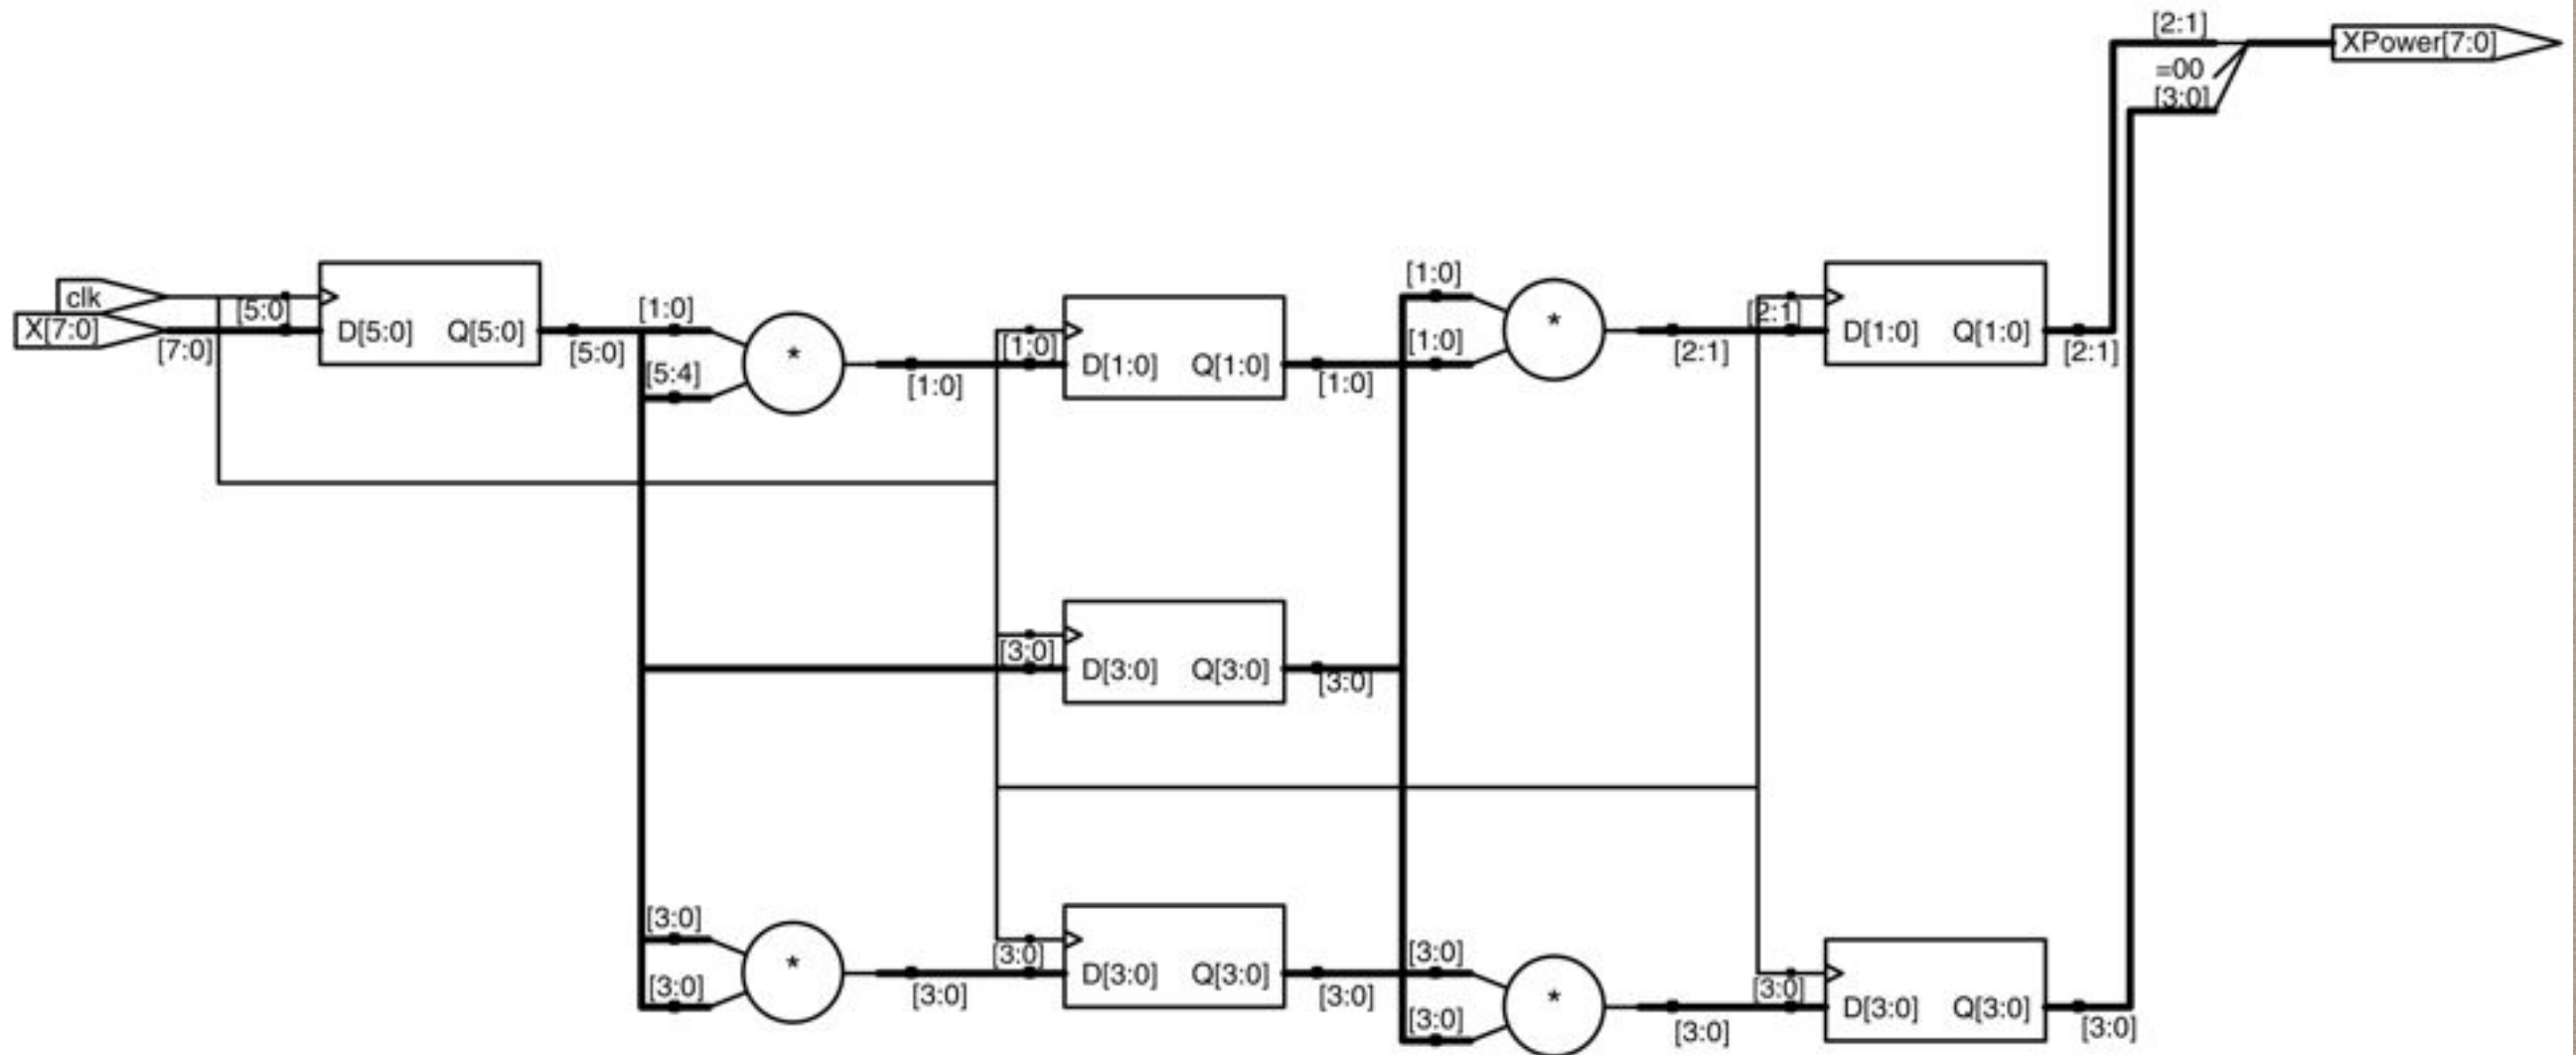


# TIMING -- PARALLEL STRUCTURES

THROUGHPUT =

LATENCY =

CRITICAL PATH =





# TIMING -- FLATTEN LOGIC STRUCTURES

DEMULTIPLEXING:  
DIRECT 'in' TO ONE OF FOUR REGISTERS  
DEPENDING UPIN 'ctrl'

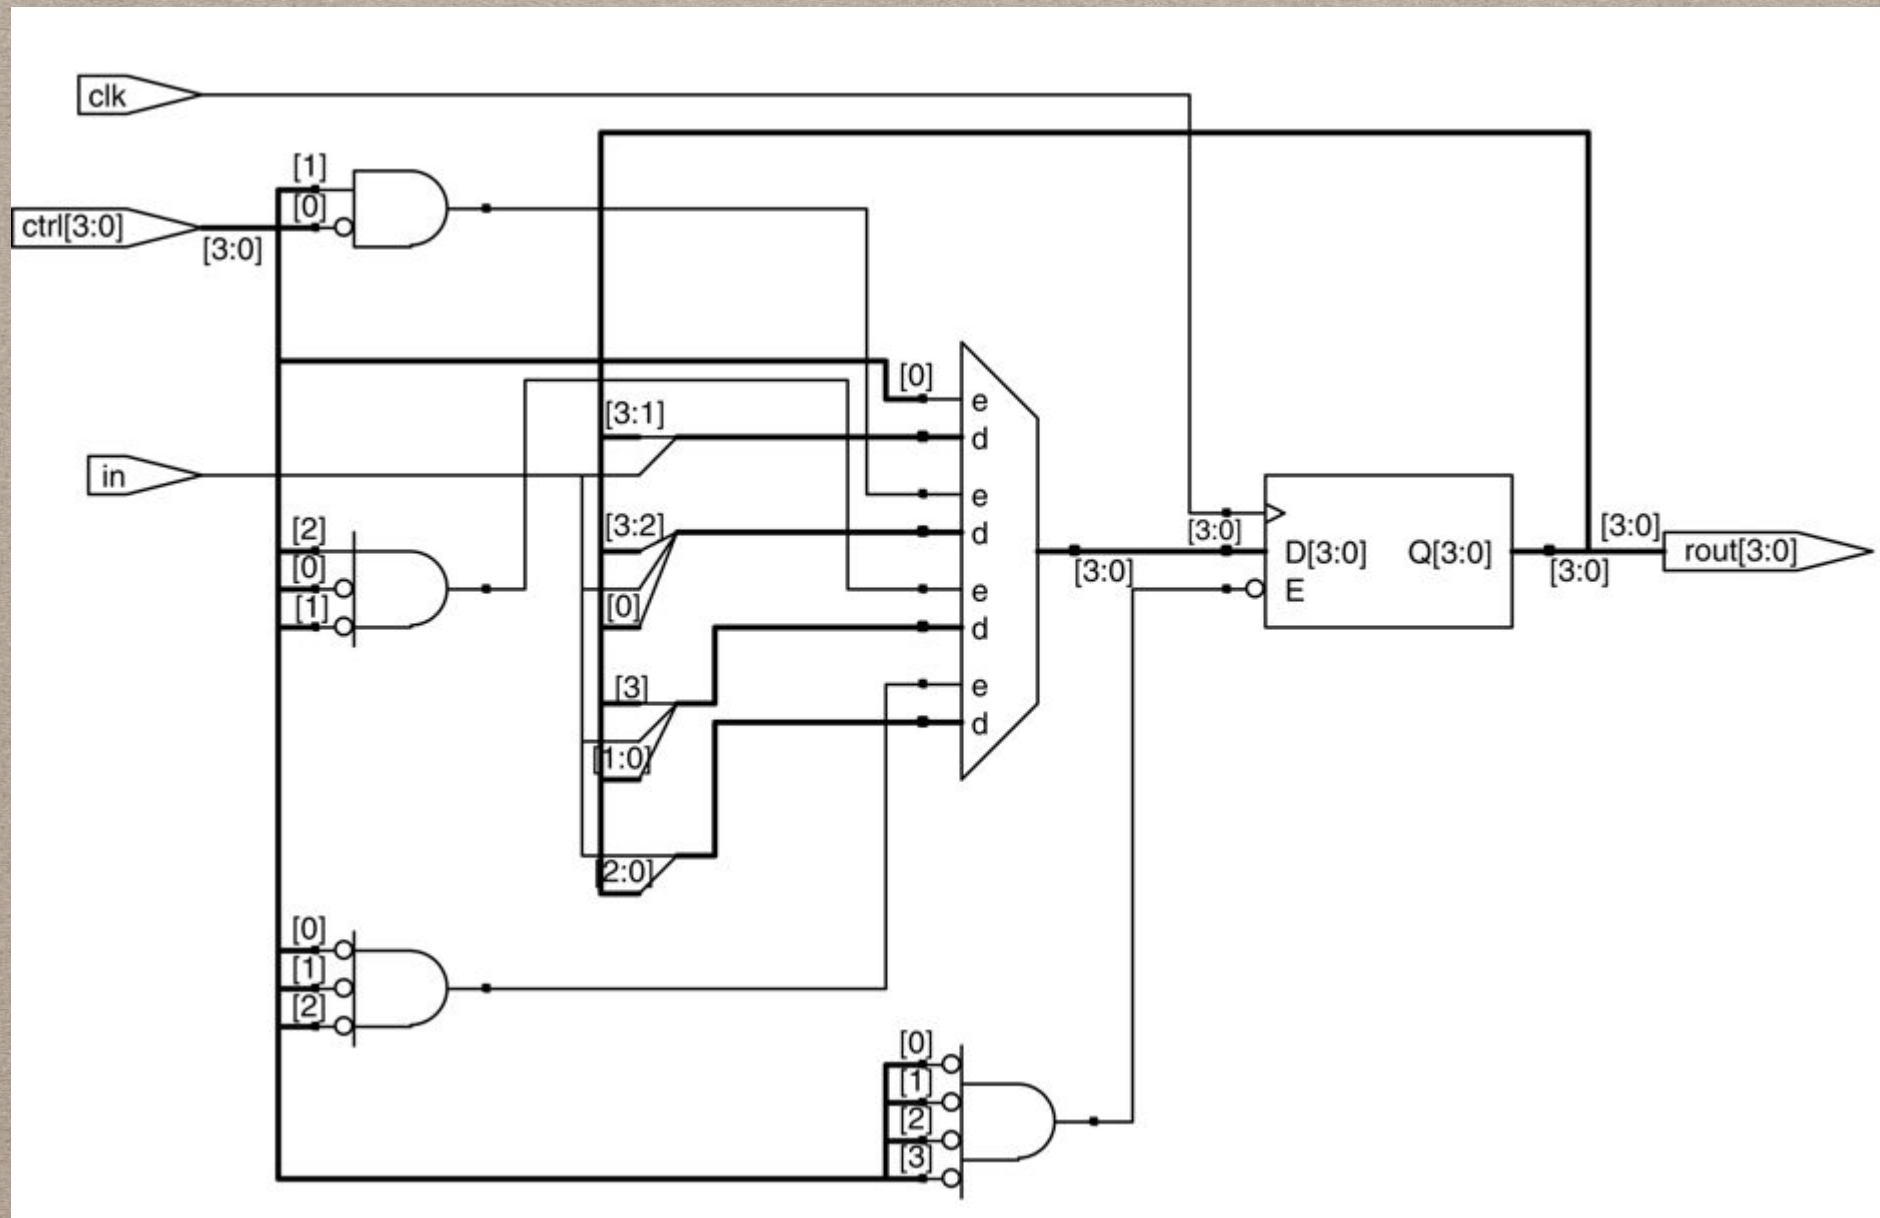
ASSUMED :  
COUT IS *ONE-HOT ENCODED*

```
module regwrite(  
    output reg [3:0] rout,  
    input            clk, in,  
    input            [3:0] ctrl);  
  
    always @(posedge clk)  
        if(ctrl[0])      rout[0] <= in;  
        else if(ctrl[1]) rout[1] <= in;  
        else if(ctrl[2]) rout[2] <= in;  
        else if(ctrl[3]) rout[3] <= in;  
  
endmodule
```

NOTE THAT THE CONDITIONALS AND OUTPUT  
ASSIGNMENTS ARE INDEPENDENT OF EACH  
OTHER FOR EACH  
IF STATEMENT



# TIMING -- FLATTEN STRUCTURES



## PRIORITY ENCODER!



# TIMING -- FLATTEN LOGIC STRUCTURES

```
module regwrite(  
    output reg [3:0] rout,  
    input            clk, in,  
    input            [3:0] ctrl);  
  
    always @(posedge clk) begin  
        if(ctrl[0]) rout[0] <= in;  
        if(ctrl[1]) rout[1] <= in;  
        if(ctrl[2]) rout[2] <= in;  
        if(ctrl[3]) rout[3] <= in;  
    end  
endmodule
```

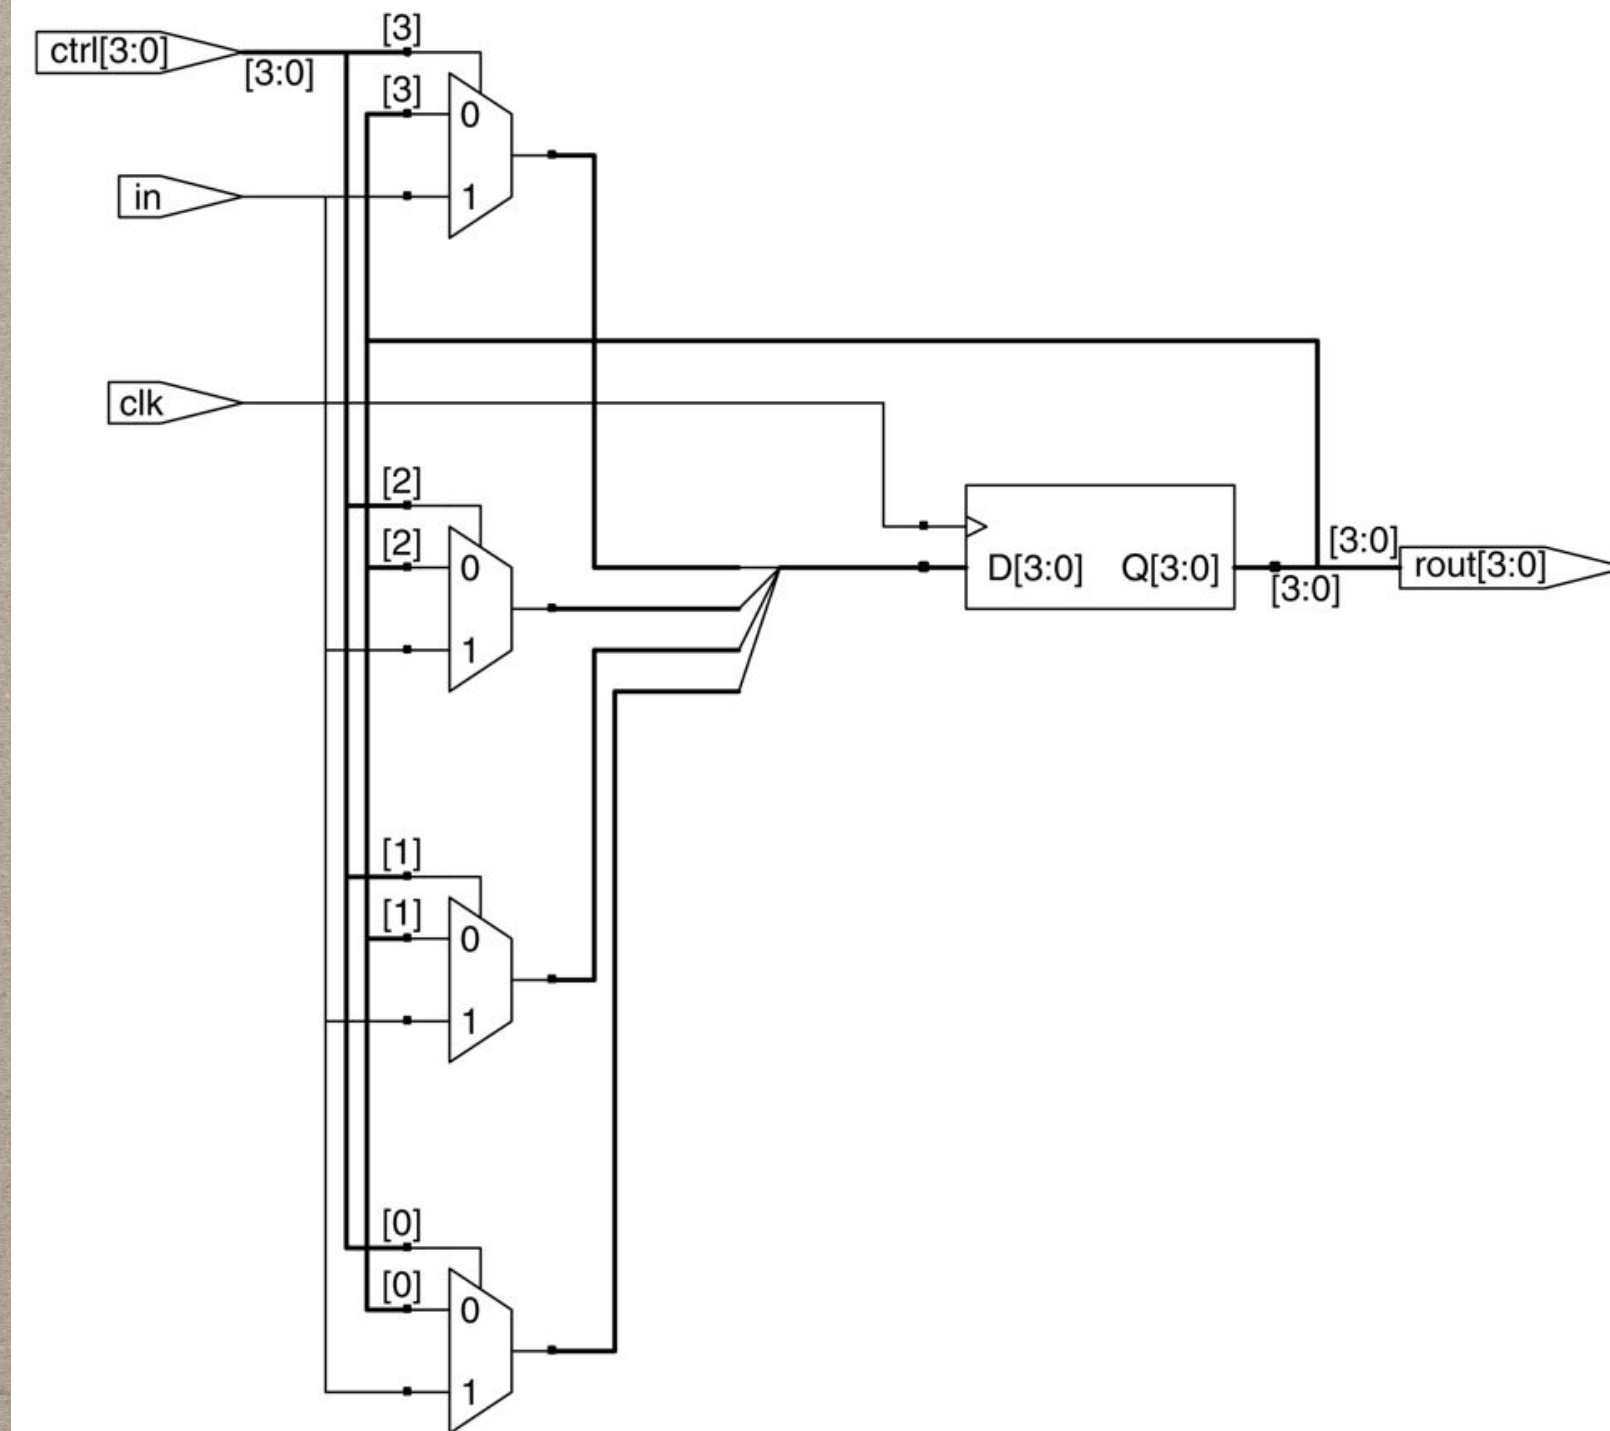


# TIMING -- FLATTEN STRUCTURES

THROUGHPUT =

LATENCY =

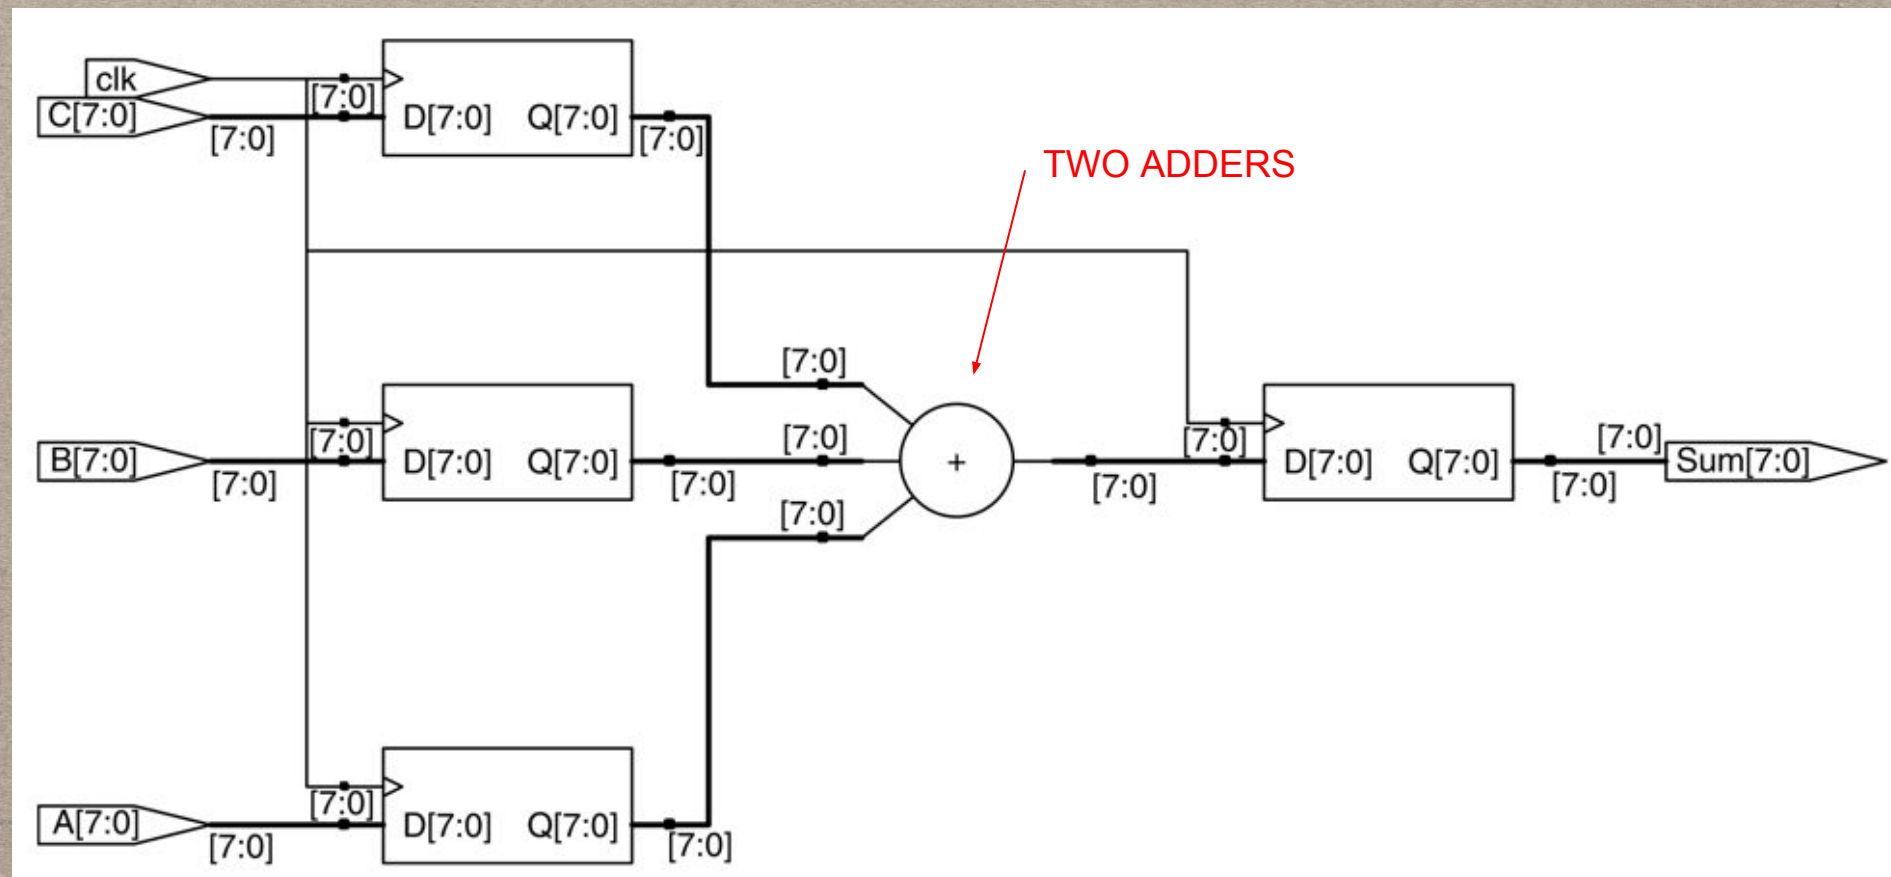
CRITICAL PATH =





# TIMING -- REGISTER BALANCING

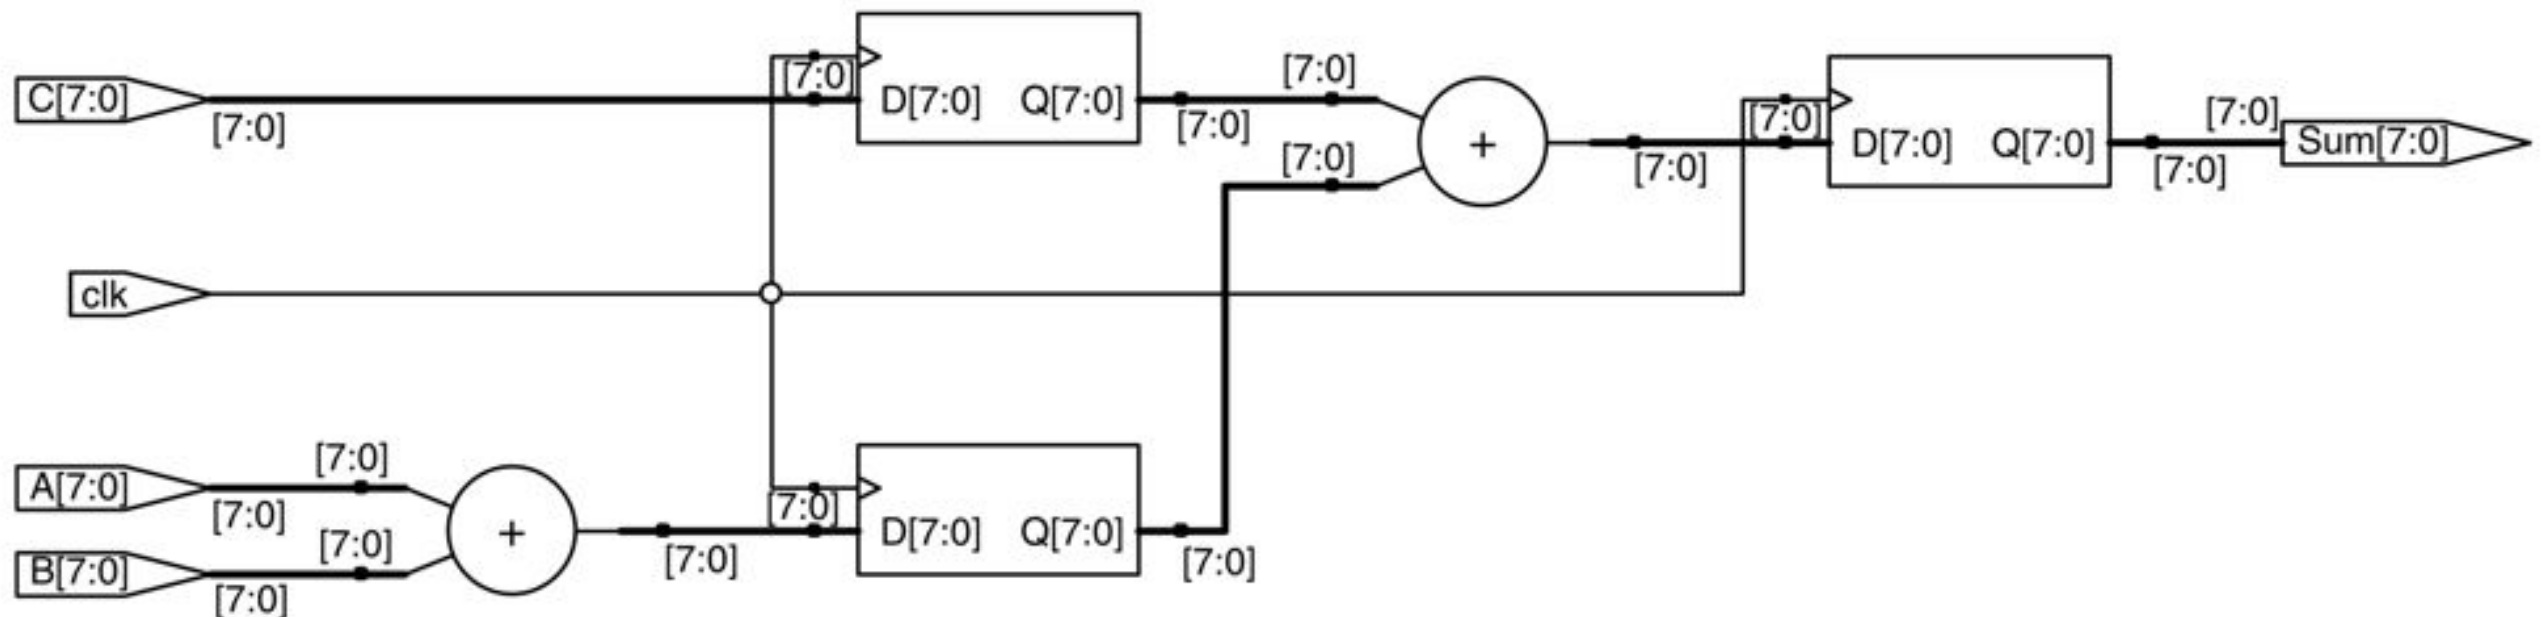
```
module adder(  
    output reg [7:0] Sum,  
    input [7:0] A, B, C,  
    input clk);  
    reg [7:0] rA, rB, rC;  
    always @(posedge clk) begin  
        rA <=A;  
        rB <=B;  
        rC <=C;  
        Sum <=rA+rB+rC;  
    end  
endmodule
```





# TIMING -- REGISTER BALANCING

```
module adder(  
    output reg [7:0] Sum,  
    input      [7:0] A, B, C,  
    input      clk);  
    reg        [7:0] rA, rB, rC;  
    always @(posedge clk) begin  
        rABSum <= A + B;  
        rC      <= C;  
        Sum     <= rABSum + rC;  
    end  
endmodule
```






# TIMING -- REORDER PATHS

```
module randomlogic(  
    output reg [7:0] Out,  
    input      [7:0] A, B, C,  
    input      clk,  
    input      Cond1, Cond2);  
always @(posedge clk)  
    if(Cond1  
        Out <= A;  
    else if(Cond2 && (C < 8))  
        Out <= B;  
    else  
        Out <= C;  
endmodule
```

THROUGHPUT =

LATENCY =

CRITICAL PATH =






# TIMING -- REORDER PATHS

```
module randomlogic(  
    output reg [7:0] Out,  
    input [7:0] A, B, C,  
    input clk,  
    input Cond1, Cond2);  
wire CondB = (Cond2 & !Cond1);  
always @(posedge clk)  
    if(CondB && (C < 8))  
        Out <= B;  
    else if(Cond1)  
        Out <= A;  
    else Out <= C;  
endmodule
```

THROUGHPUT =

LATENCY =

CRITICAL PATH =





THE END