


# CHAPTER 3 LOW-POWER DESIGN

## ECE4514


### DIGITAL DESIGN 2



# Topics


- Sources of power consumption
- The impact of clock control on dynamic power consumption
- Problems with clock gating
- Managing clock skew on gated clocks
- Input control for power minimization
- Impact of the core voltage supply
- Guidelines for dual-edge triggered flip-flops

# Power in the 20th Century



Power increases despite Vdd decrease

# Power Density




# Power Dissipation

- Dynamic power
- Leakage power
- Glitches
- Short circuit power

# Dynamic Power

- Occurs at each switching
- $P_d = C_L \bullet V_{dd}^2 \bullet f_p / 2$
- $f_p$  switching frequency



AN INVERTER SWITCHING

# Dynamic Power

When signals wiggle  
power is consumed

# Reduce Dynamic Power

$$P_d = 1/2 [ \alpha \bullet C_L \bullet V_{dd}^2 \bullet f_p ]$$

- $\alpha$ : *Activity Factor*  
clock gating, sleep mode
- $C$ : use small transistors (esp. on clock), short wires
- $V_{DD}$ : lowest suitable voltage
- $f$ : lowest suitable frequency




# Leakage Current



Main Sources of Leakage	Impact	Mitigation Techniques
Subthreshold leakage ( $I_{\text{sub}}$ )	Dominant	<ul style="list-style-type: none"><li>■ Lower voltage</li><li>■ Higher voltage threshold</li><li>■ Longer gate length</li><li>■ Dopant profile optimization</li></ul>
Gate direct-tunneling leakage ( $I_{\text{G}}$ )	Dominant	High-k metal gate (HKMG)
Gate-induced gate leakage ( $I_{\text{GIDL}}$ )	Small	Dopant profile optimization
Reverse-biased junction leakage current ( $I_{\text{REV}}$ )	Negligible	Dopant profile optimization

# Subthreshold Leakage Current

- Killer to CMOS technology




# Leakage Current




Where are the diodes?

# Trends in Mobile Devices




# Moore and Moore Transistors




# Leakage Trends

Source: ITRS Roadmap 2007





# I/O Power




Main Factors Impacting I/O Power	Mitigation Techniques
Termination resistors (on-chip series termination ( $R_S$ OCT) and on-chip parallel termination ( $R_T$ OCT))	Dynamic on-chip termination (DOCT)
Output buffer drive strength	Programmable drive strength
Output buffer slew rate	Programmable slew rate
I/O standard (single ended, voltage referenced, or differential)	Support for multiple I/O standards
Voltage supply	Support for various voltage rails
Capacitive load (charging/discharging)	Interface dependent

# Glitches




State: 0 1 2 3 4




Decoding Glitch Example



# Glitch Removal




(a) Original circuit with glitch



(b) Glitch removed by delaying input c

# Short Circuit Power

- During switching, there is a short moment when both PMOS and CMOS are partially on
- $P_s = Q \bullet (V_{dd} - V_t)^3 \bullet t_r \bullet f_p$
- $t_r$  rising time



# Power Mitigation

Things we can do as logic designers


# Clock Gating

- Gate off clock to idle functional units
  - e.g., floating point units
  - Need logic to generate **disable** signal
    - increases complexity of control logic
    - timing critical to avoid clock glitches
  - Additional gate delay on clock signal
    - gating gate can replace a buffer in the clock distribution tree

# Clock Gating

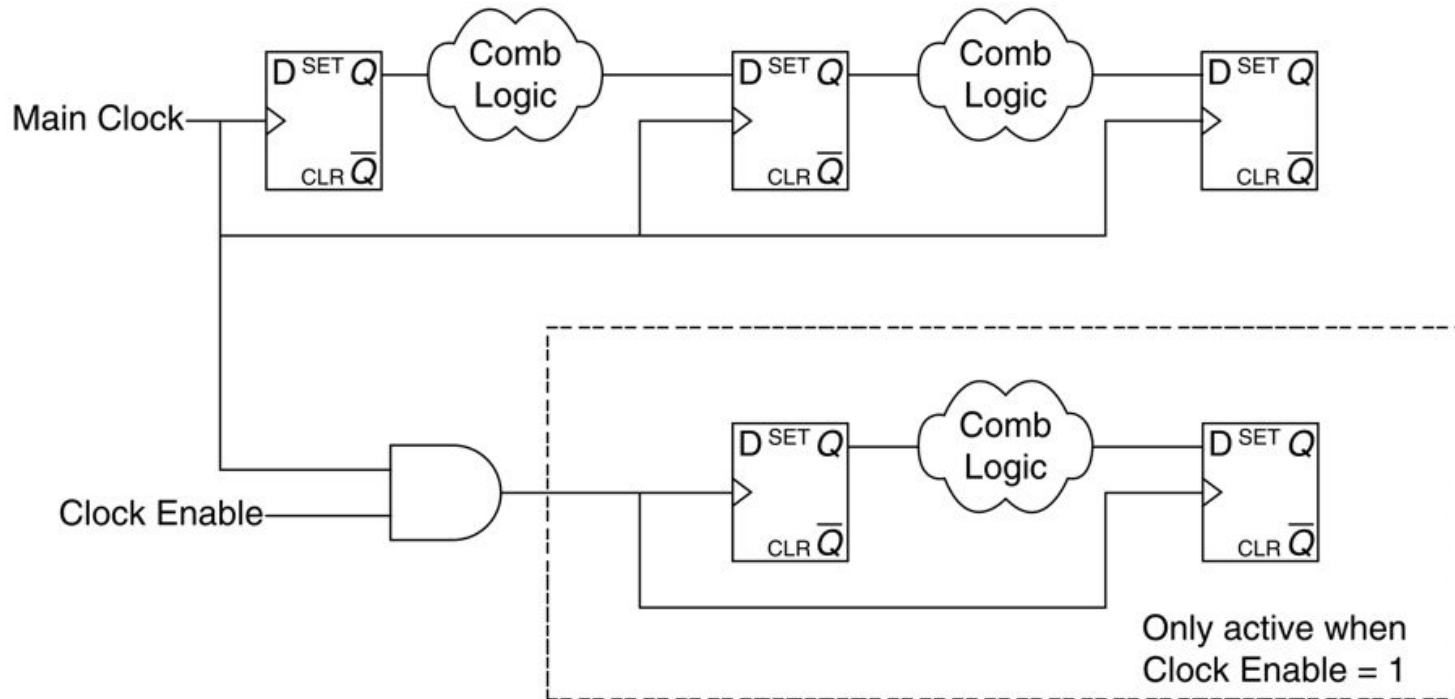
A gated clock introduces a new clock domain and will create difficulties for the designer.

Figure 1-5 Clock Gating Example




generally  
viewed as a BAD  
THING TO DO.

# Clock Gating



# Clock Gating vs Clock Enable



No impact on timing, yet,  
clock distribution network  
still wiggling.

```
always@(posedge clock) begin
    if (in1 & in2 & load) begin
        out1 <= data;
    end
end
```

# Clock Skew Danger

Mishandling clock skew can cause  
“catastrophic failures” in the FPGA.



# Simplistic Clock Gating

```
// Poor design practice
module clockgating(
    output dataout,
    input  clk, datain,
    input  clockgate1);
    reg    ff0, ff1, ff2;
    wire   clk1;
```

```
// clocks are disabled when gate is low
```



```
    assign clk1      = clk & clockgate1;
    assign dataout = ff2;
```

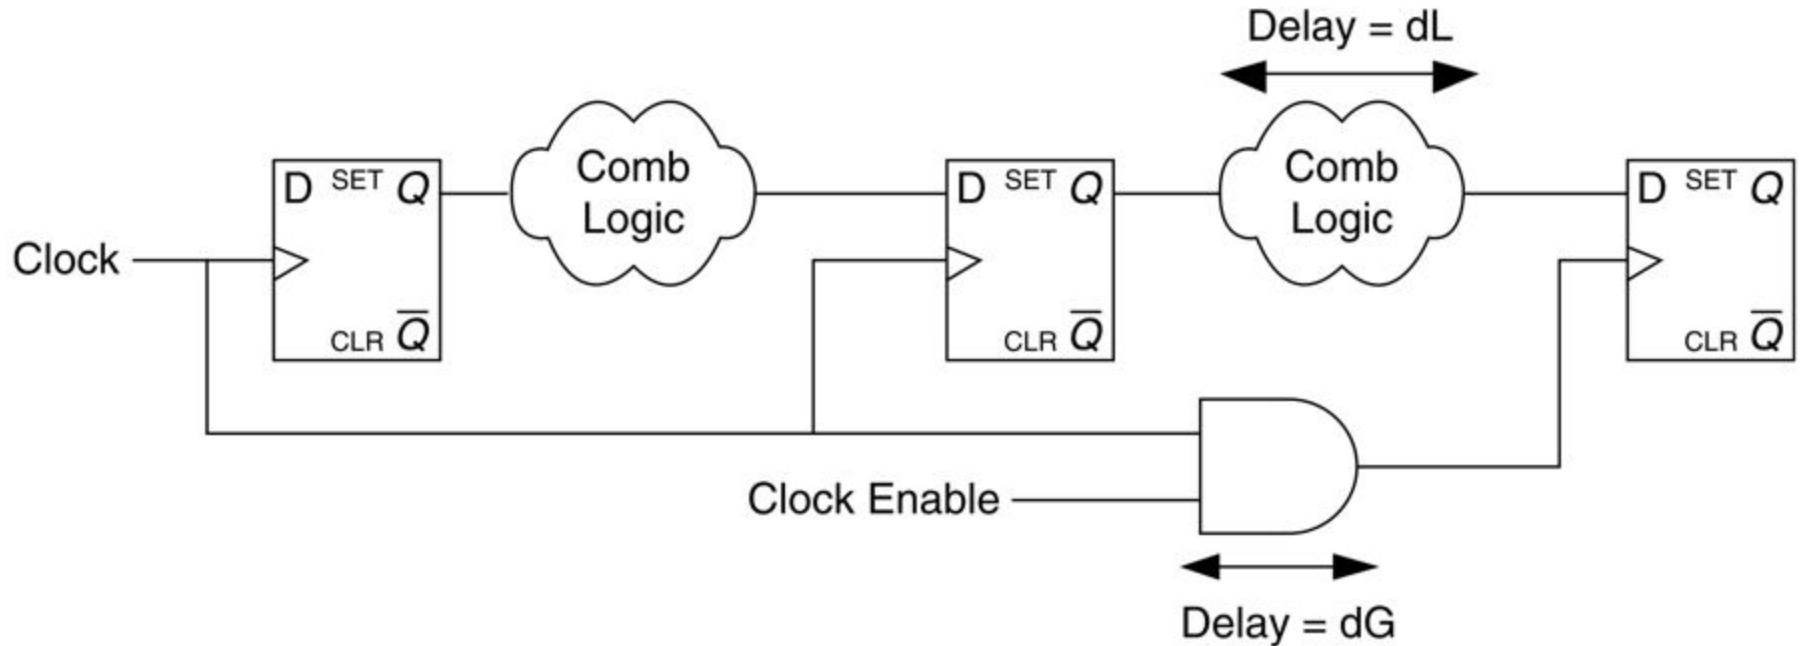
```
    always @(posedge clk)
        ff0 <= datain;
```

```
    always @(posedge clk)
        ff1 <= ff0;
```




```
    always @(posedge clk1)
        ff2 <= ff1;
endmodule
```

# Poor Design Practice




# Timing

$$F_{\max} = \frac{1}{T_{\text{clk-q}} + T_{\text{logic}} + T_{\text{routing}} + T_{\text{setup}} - T_{\text{skew}}}$$




# Implemented Design



# Alternative Design

```
module noclockgating(  
    output dataout,  
    input clk, datain, input clockgate1);  
    reg ff0, ff1, ff2; wire clk1;  
    // clocks are disabled when gate is  
    low  
    assign dataout = ff2;  
    always @(posedge clk)  
        ff0 <= datain;  
    always @(posedge clk)  
        ff1 <= ff0;  
    always @(posedge clk)  
        if (clockgate1) ff2 <= ff1;  
endmodule
```

*REDO THE  
DESIGN SO  
THAT IT USES A  
CLOCK ENABLE  
INSTEAD*



# Poor Design Practice

```
wire [8:0] count, cntout;
```

```
prescale pre( .clk(clk), .rst(rst),  
              .count(count) );
```

```
maincount main ( .clk(count[8]), .ce(1),  
                 .mainout( cntout ) );
```

# Better Design Practice

```
wire [8:0] count, cntout;
```

```
prescale pre( .clk(clk), .rst(rst),  
  .count(count) );
```

```
assign ce = (count = 9'd38);
```

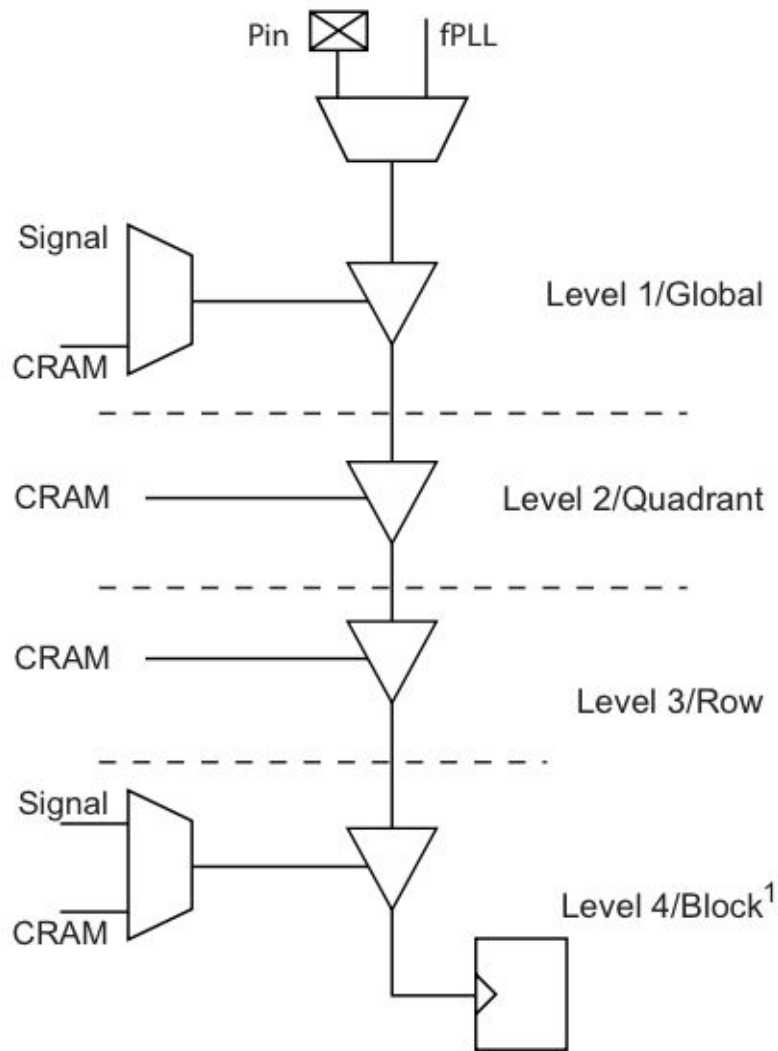
```
maincount main (.clk(clk), .ce(ce),  
  .mainout( cntout ));
```

# Stratix Clock Gating



- **Static clock gating**—The clock is enabled or disabled using a configuration RAM bit at programming time. The Quartus II software leverages static clock gating on clocks that are either always used or never used in the design.
- **Dynamic clock gating**—The clock is enabled or disabled using a signal generated by designers or by the Quartus II software. Quartus II software synthesizes designs and generates signals at various clock stages to dynamically shut off clocks when the downstream logic does not need to toggle.



# Stratix Clock Gating



# Controlling Rise/Fall Time



# Controlling Rise/Fall Time

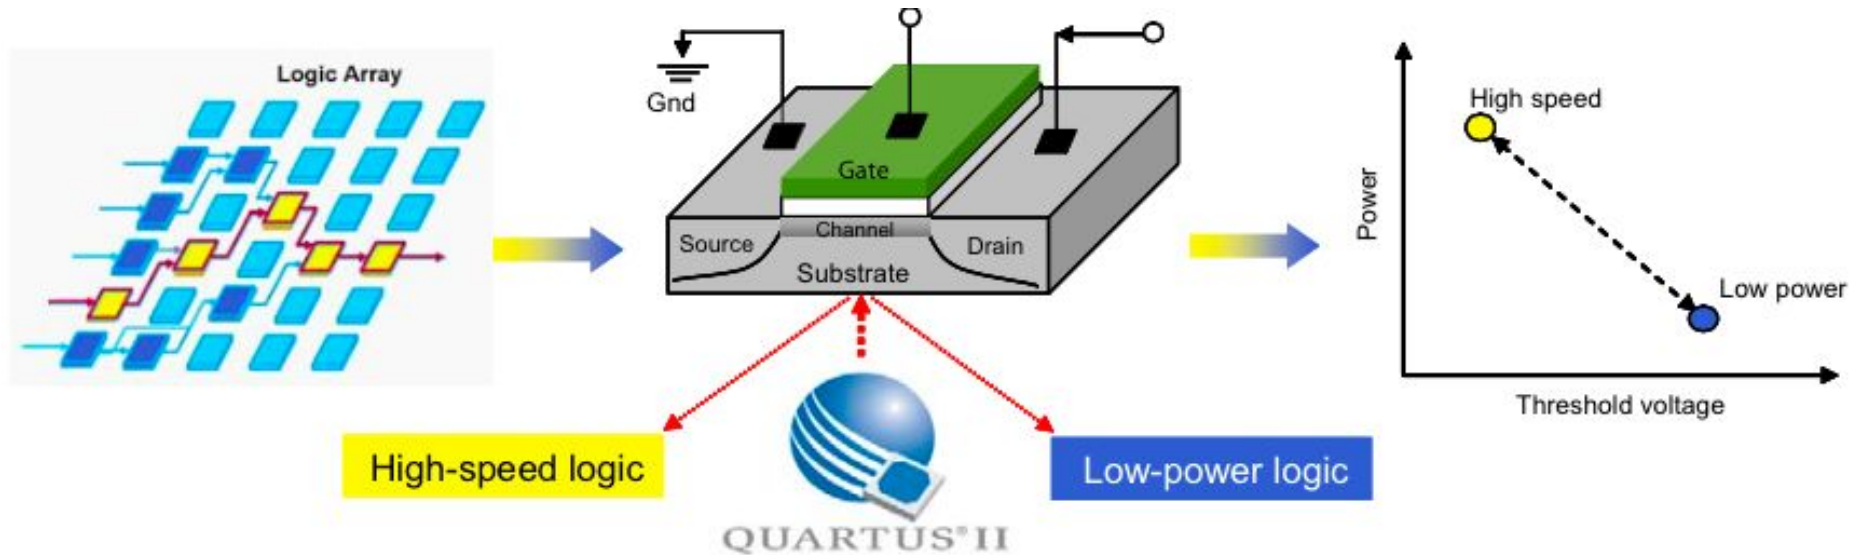
To minimize the power dissipation of input devices, minimize the rise and fall times of the signals that drive the input.

What about outputs?

# Supply Voltage

Dynamic power dissipation drops off with the square of the core voltage, but reducing voltage will have a negative impact on performance.

# Altera Programmable Power



# Altera Programmable Power

- **Low-power mode**—Quartus II software reduces the back bias voltage (making it more negative), which increases the threshold voltage of transistors. This reduction minimizes sub-threshold leakage currents and unwanted static power in non-timing-critical circuit paths.
- **High-performance mode**—Quartus II software increases the back bias voltage (making it less negative), which reduces the threshold voltage of transistors for faster switching. This performance increase is required on logic in the few timing-critical paths to help meet the design's specified timing constraints and deliver maximum performance.

# Dynamic Voltage and Frequency Scaling

- Run at the lowest supply voltage that meets the timing constraints
  - DFS (dynamic frequency scaling)
  - DVS (dynamic voltage scaling)
- A DVS+DFS system requires the following
  - A programmable clock generator (PLL)
    - PLL from 200MHz  $\rightarrow$  700MHz in increments of 33MHz
  - A supply regulation loop that sets the minimum  $V_{DD}$  necessary for operation at the desired frequency
  - An operating system that sets the required frequency + supply voltage to meet the task completion deadlines

# Dual Edge Flip-Flops

```
module duedge(  
  output reg dataout, input clk, datain);  
  reg ff0, ff1;  
  always @(posedge clk)  
    ff0    <= datain;  
  always @(posedge clk or negedge clk) begin  
    ff1 <= ff0;  
    dataout <= ff1;  
  end  
endmodule
```




# Dual Edge Flip-Flops


Dual-edge triggered flip-flops reduce the energy in the clock distribution network by 50%.

Dual-edge triggered flip-flops should only be used if they are provided as primitive elements (and they are NOT in Altera devices)

# Power Consumed



# FPGA Power Consumption



Most FPGA power consumption is attributed to losses in the programmable interconnect

# Altera Gated Clocks

**Entering the Danger Zone**

# FPGA Gated Clocks

In a design, you can use combinational logic as a clock signal that can prevent certain logic in a circuit from being activated by the clock signal and can therefore reduce the total power consumption of a device. The combinational logic that is used as a clock signal should follow the following guidelines


# FPGA Gated Clocks

1. The gating logic should be a two-input AND gate or a two-input OR gate.
2. There should be only one input pin that acts as a primary input clock signal to the AND or gate.
3. The non-clock input clock enable signal to the AND or OR gate should be synchronized with a register. The register should be clocked by the same input pin that acts as the primary input clock signal to the AND or OR gate.
4. If the combinational logic uses an AND gate, the clock port of the register that drives the AND gate should be active on the falling edge and the clock port of the register driven by the AND gate should be active on the rising edge.
5. *or*

# FPGA Gated Clocks

6. If the combinational logic uses an OR gate, the clock port of the register that drives the OR gate should be active on the rising edge and the clock port of the register driven by the OR gate should be active on the falling edge.
7. To reduce the probability of glitches, the `AND` gate should be placed close to the flipflop that generates the gating signal.
8. When specifying timing constraints, Altera recommends enclosing the gating circuitry (`AND` gate and flipflop) in a module. This allows you to specify timing constraints at the module output.

# FPGA Gated Clock



NOTE: TWO CLOCK DOMAINS!!!!



# Quartus Gated Clocks

## Design Assistant

Specify the potential design problems that you want the Design Assistant to check. You can choose to check the design for individual problems, or a category of design problems.

☒ Run Design Assistant during compilation

Select the rules you want the Design Assistant to apply to the project:

- ☒ Design Assistant configuration rule names
  - ☒ Clock
    - ☒ Rule C101: Gated clock should be implemented according to the Altera standard scheme
    - ☒ Rule C102: Logic cell should not be used to generate an inverted clock signal
    - ☒ Rule C103: Gated clock is not feeding at least a pre-defined number of clock port to effectively save power : 30
    - ☒ Rule C104: Clock signal source should drive only clock input ports
    - ☒ Rule C105: Clock signal should be a global signal : 25
    - ☒ Rule C106: Clock signal source should not drive registers triggered by different clock edges
  - ☒ Reset
  - ☒ Timing closure
  - ☒ Non-synchronous design structure
  - ☒ Signal race
  - ☒ Asynchronous clock domains

# **Design Techniques to Reduce Power from Glitches: Case Study**

## **Grey Encoding for Counters**

Study by Grant Lewis

# Binary Encoding

- 3-bit Counter has 14 transitions in 8 cycles
  - $14/8 = 1.75$  transitions per cycle
- N-bit Counter has  $2(2^N - 1)$  transitions
  - Transitions per cycle  $\rightarrow 2$  as  $N \rightarrow \infty$

4-bit	1.875 transitions/clock
5-bit	1.9375
6-bit	1.96875
7-bit	1.984375
8-bit	1.9921875
9-bit	1.99609375

000	--
001	1
010	2
011	1
100	3
101	1
110	2
111	1
000	3

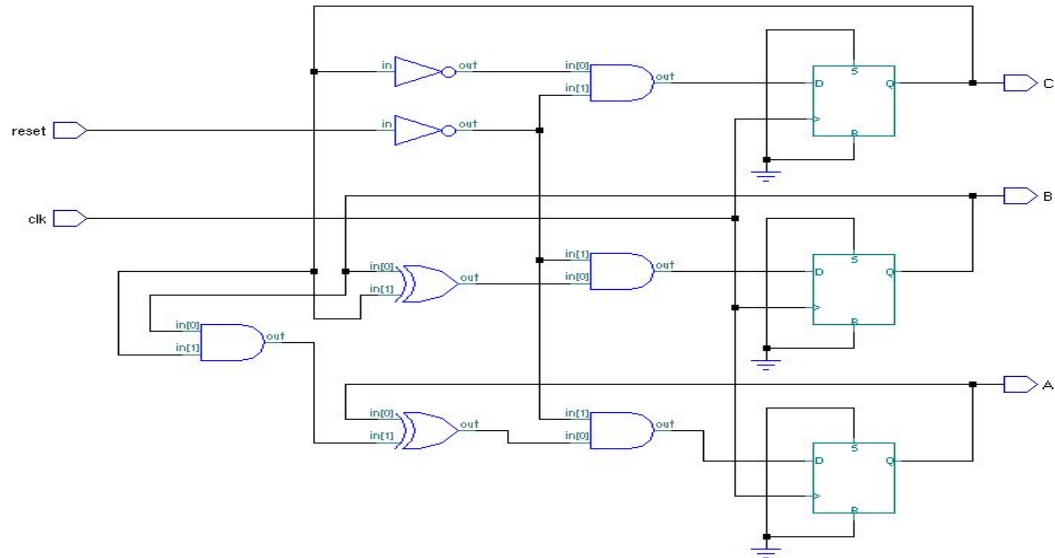
# Binary Encoding: 3-bit

- $A = a \oplus bc$

- $B = b \oplus c$

- $C = \overline{c}$

- 8 logic gates



Synthesized with Leonardo Spectrum → Standard Cell

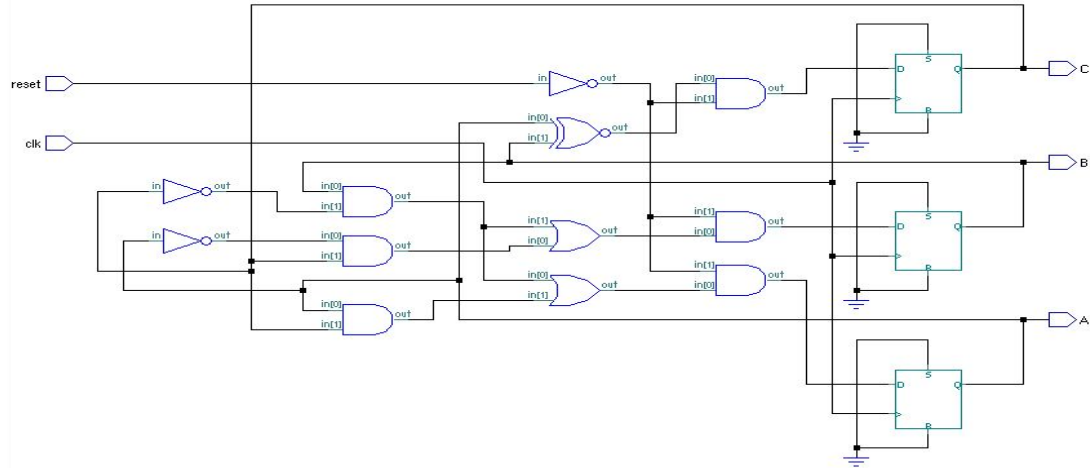
# Gray Encoding

- 3-bit counter has 8 transitions in 8 cycles
  - $8/8 = 1$  transition per cycle
- N-bit counter has  $2^N$  transitions
  - Transitions per cycle is 1 for any size counter

000	--
001	1
011	1
010	1
110	1
111	1
101	1
100	1
000	1

# Grey Encoding

- $A = \overline{bc} + ac$
- $B = \overline{ac} + \overline{bc}$
- $C = \overline{a \oplus b}$



- Adds 4 logic gates (total 12 gates)

# Power Analysis

(0.18  $\mu\text{m}$ , 1.8V supply)

	Decimal Encoding	Gray Encoding	% Reduction
<b># of Gates</b>	8	12	-50%
<b># of Glitches</b>	3122	1618	48.17%
<b>Glitch Power (<math>\mu\text{W}</math>)</b>	1.852307	0.934297	49.56%
<b>Dynamic Power (<math>\mu\text{W}</math>)</b>	1.855127	0.940320	49.31%
<b>Leakage Power (<math>\mu\text{W}</math>)</b>	.914450	.889597	2.72%
<b>Total Power (<math>\mu\text{W}</math>)</b>	2.769578	1.829918	<b>34.18%</b>

# Conclusion

- Much of dynamic power consumption comes from glitches
  - Can be reduced by path balancing
- Overall, gray counters are more power efficient



# Minimizing Transitions on a Bus

```
// Code that resets the bus to default  
status after valid gets de-asserted
```

```
always@(posedge clk or negedge reset)
```

```
begin  
    if(!reset)  
        data_bus <= 16'b0 ;  
    else if (data_bus_valid)  
        data_bus <= data_o ;  
    else  
        data_bus <= 16'b0 ;  
end
```

```
// Code that holds the bus to its previous  
value after valid gets de-asserted
```

```
always@(posedge clk or negedge reset)
```

```
begin  
    if(!reset)  
        data_bus <= 16'b0 ;  
    else if (data_bus_valid)  
        data_bus <= data_o ;  
end
```

**The End**

# Leakage Trends

